

Getting Started with Microservices and Kubernetes: Overview

What We'll Learn?

Section 1

Getting Started with
Kubernetes

Section 3

Deploying Stateful Services

Section 5

Continuous Delivery with
Kubernetes

Section 2

The Kubernetes Deployment Model

Section 4

Managing Configuration Data

Prerequisites

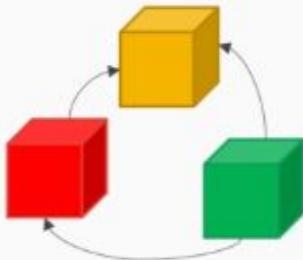
What do you need to know?

docker

Working knowledge
of Docker

Prerequisites

What do you need to know?



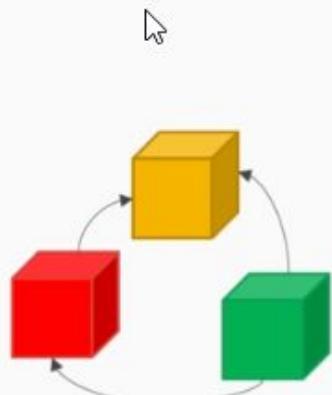
Microservice
basics



Working knowledge
of Docker

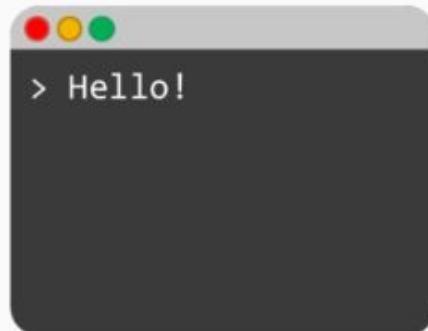
Prerequisites

What do you need to know?



Microservice
basics

docker



Working knowledge
of Docker

Some command-
line skills

Required Software

(Don't Worry, It's All Free!)

kubectl

Command-Line Client for Kubernetes

- Needs to run on your local machine
- Installation instructions:
<https://kubernetes.io/docs/tasks/tools/install-kubectl/>
- Varies by operating system

Minikube

Kubernetes Running Locally

- Download from:
<https://github.com/kubernetes/minikube/releases>
- Requires Oracle VirtualBox
<https://www.virtualbox.org/wiki/Downloads>
- Alternative: Running Kubernetes cluster in cloud, for example, GKE or AKS

Your Favorite Code Editor

Vim or Emacs?

- Recommendation for the indecisive: Visual Studio Code
- [https://code.visualstudio.com/
download](https://code.visualstudio.com/download)

Setting up with Minikube

Prerequisite – Oracle VirtualBox

- Download at:
<https://www.virtualbox.org/wiki/Downloads>



Minikube Download

<https://github.com/kubernetes/minikube/releases>

 [kubernetes / minikube](#)

[Watch](#) 315 [Unstar](#) 8,716 [Fork](#) 1,214

[Code](#) [Issues 292](#) [Pull requests 27](#) [Projects 1](#) [Wiki](#) [Insights](#)

[Releases](#) [Tags](#)

[Latest release](#)

 v0.26.1
 minikube-bot released this 16 days ago · 17 commits to master since this release
 6ded2b6

Assets

 docker-machine-driver-hyperkit	25.6 MB
 docker-machine-driver-kvm2	35.4 MB
 localkube	173 MB
 localkube.sha256	65 Bytes
 minikube-darwin-amd64	40.6 MB
 minikube-darwin-amd64.sha256	65 Bytes

↳ v0.25.2
→ 67a6b82

v0.25.2

💻 minikube-bot released this on 22 Mar · 149 commits to master since this release

Assets

📄 docker-machine-driver-hyperkit	26.2 MB
📄 docker-machine-driver-kvm2	36.2 MB
📄 localkube	163 MB
📄 localkube.sha256	65 Bytes
📄 minikube-darwin-amd64	41.3 MB
📄 minikube-darwin-amd64.sha256	65 Bytes
📄 minikube-installer.exe	11.1 MB
📄 minikube-linux-amd64	41.5 MB
📄 minikube-linux-amd64.sha256	65 Bytes
📄 minikube-windows-amd64	41.5 MB
📄 minikube 0.25-2.deb	7.37 MB

```
chmod +x Downloads/minikube-darwin-amd64
sudo mv Downloads/minikube-darwin-amd64 /usr/local/bin/minikube
minikube status
minikube:
cluster:
kubectl:
✗ ➔ minikube start
Starting local Kubernetes v1.9.4 cluster...
Starting VM...
Getting VM IP address...
Moving files into cluster...
Setting up certs...
Connecting to cluster...
Setting up kubeconfig...
Starting cluster components...
Kubectl is now configured to use the cluster.
Loading cached images from config file.
➔ minikube stop
Stopping local Kubernetes cluster...
Machine stopped.
➔ minikube delete
```

Minikube Add-Ons

- Enabled per default:
 - DNS
 - Dashboard
 - Storage provisioner

Minikube Add-Ons

- Enabled per default:
 - DNS
 - Dashboard
 - Storage provisioner
- Useful to enable:
 - Ingress
 - Heapster

```
➔ minikube addons enable ingress  
ingress was successfully enabled  
➔ minikube addons enable heapster  
heapster was successfully enabled  
➔
```

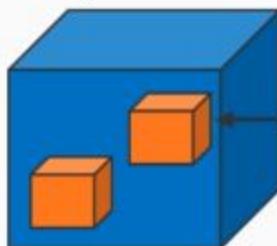
Deploying a Microservice

Kubernetes Core Concepts

Pods – The Basic Deployment Unit

Pod

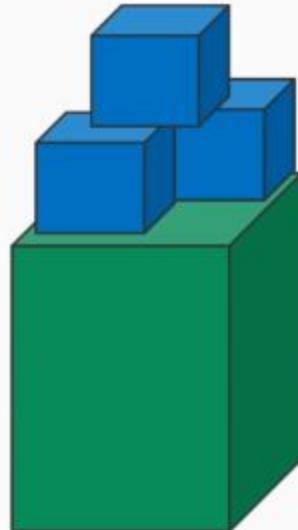
- Multiple containers (in practice seldom more than 1 or 2)
- Scheduled on **one** node



Container

- (Mostly) regular Docker (or rkt or CRIOS) container
- Shared network namespace for all containers in Pod

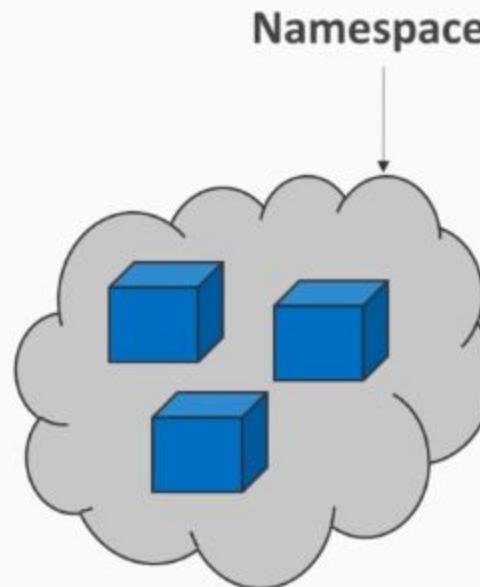
Nodes – It's Just Servers, After All



Node

- Individual server (virtualized or bare metal)
- Runs Docker engine
- Pods/containers started by Kubernetes

Namespaces – Resource Grouping and Isolation

- 
- Groups resources (like Pods)
 - In a fresh Kubernetes cluster: default and kube-system
 - Purely virtual
 - No additional isolation at container level

Defining a Pod

Assumptions

- You have a container image for your microservice available
(In this example, we will use the nginx image as a dummy image)
- Your Kubernetes cluster is authorized to pull that image
(Not an issue with the public nginx image; use imagePullSecret option otherwise)

```
~ ➔ kubectl get pods  
No resources found.
```



Pod Example

- Full documentation:
<https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod    Pod name
  namespace: default
spec:
  containers:
    - name: service
      image: nginx
      ports:
        - containerPort: 80
      env:
        - name: SOME_ENV_VAR
          value: Hello World
```

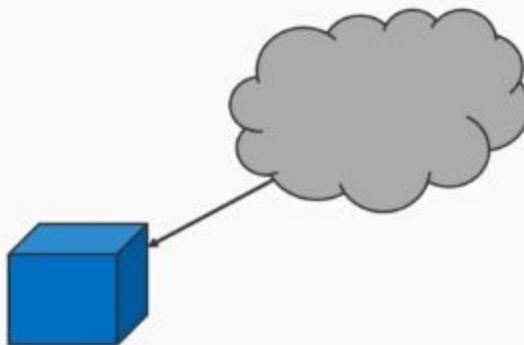
Pod Example

- Full documentation:
<https://kubernetes.io/docs/concepts/workloads/pods/pod-overview/>

```
apiVersion: v1
kind: Pod
metadata:
  name: test-pod
  namespace: default
spec:           Container specification
  containers:
    - name: service
      image: nginx
      ports:
        - containerPort: 80
      env:
        - name: SOME_ENV_VAR
          value: Hello World
```

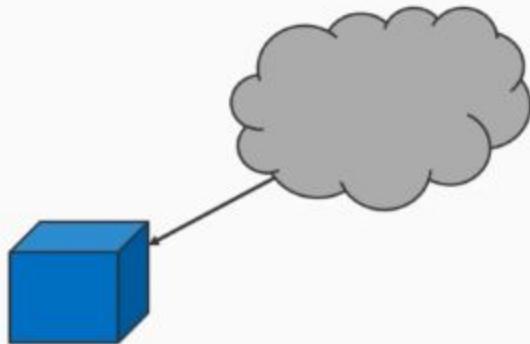
```
→ kubectl get pods
No resources found.
→ kubectl apply -f pod.yaml
pod "test-pod" created
→ kubectl get pods
NAME      READY     STATUS    RESTARTS   AGE
test-pod  1/1      Running   0          17s
→ |
```

Open Issues

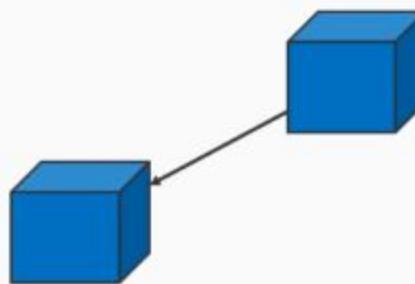


Network connectivity
from external sources

Open Issues

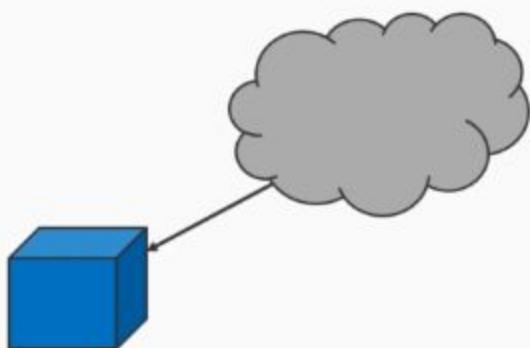


Network connectivity
from external sources

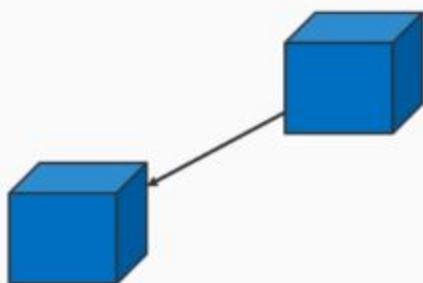


Network connectivity
to/from other services

Open Issues



Network connectivity
from external sources

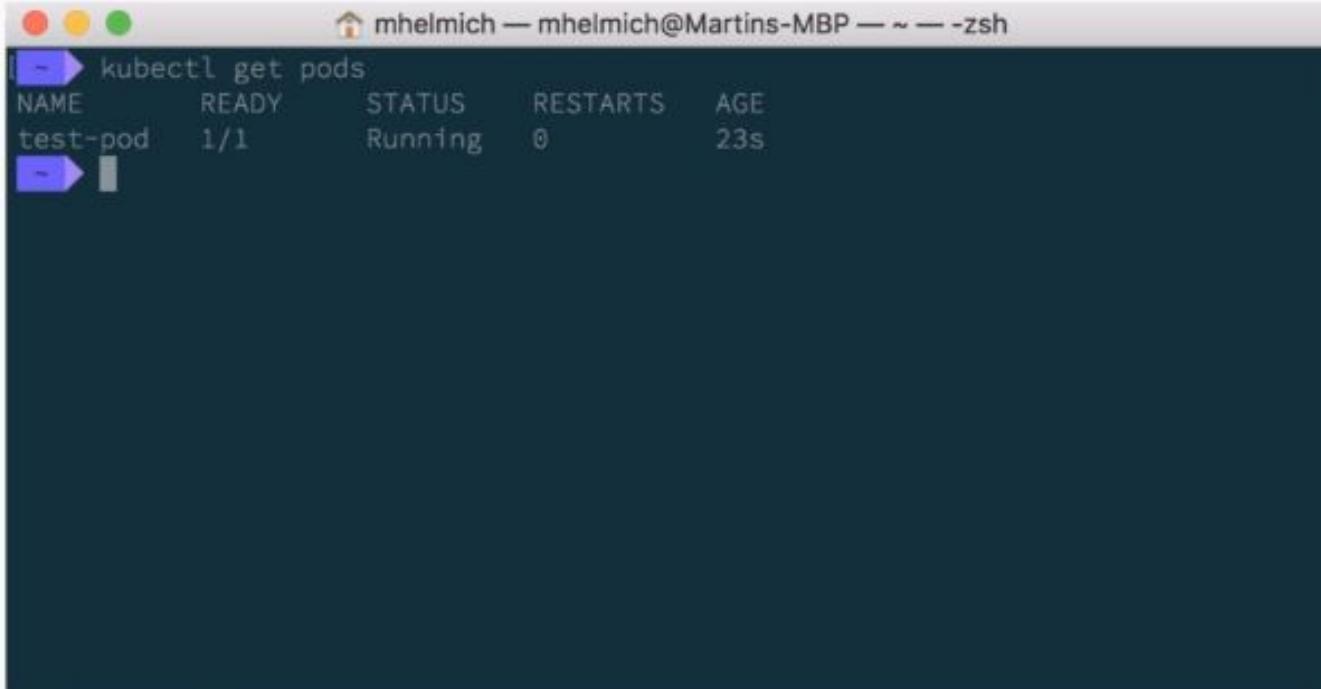


Network connectivity
to/from other services

Solution
Service and Ingress
resources, covered in
Section 2 of this course

Service Resiliency and Scalability

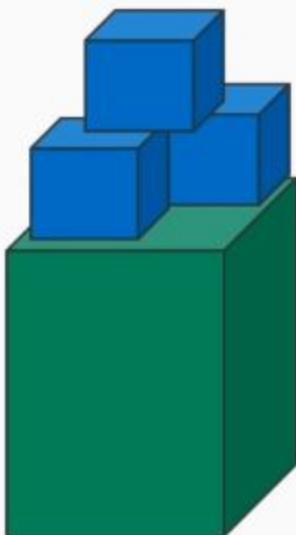
What Happened So Far?



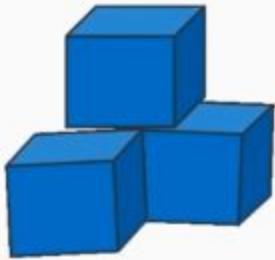
A screenshot of a macOS terminal window. The title bar shows the user's name and host as "mhelmich — mhelmich@Martins-MBP — ~ — -zsh". The command entered is "kubectl get pods". The output shows a single pod named "test-pod" with a status of "Running", one container ready, zero restarts, and an age of 23 seconds.

```
mhelmich — mhelmich@Martins-MBP — ~ — -zsh
~ ➤ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
test-pod  1/1     Running   0          23s
~ ➤
```

Pod Limitations



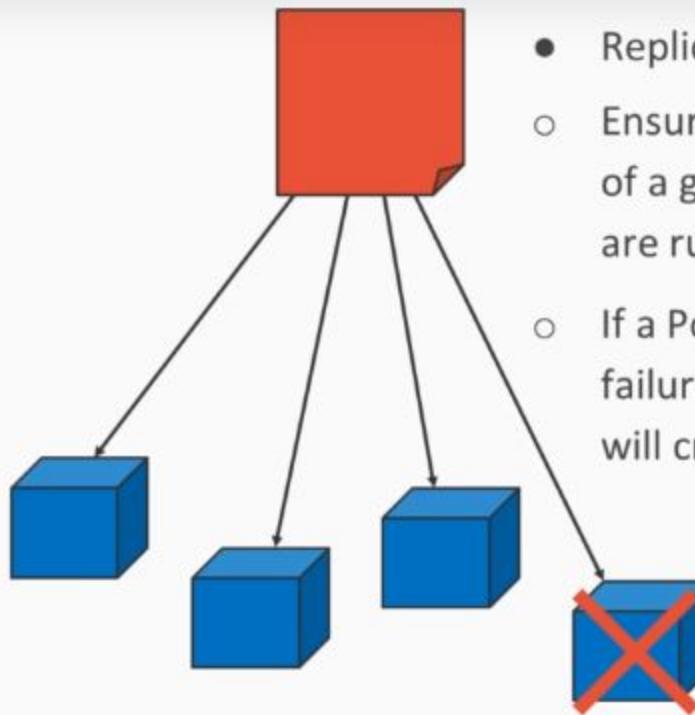
- Each Pod is assigned to a specific node in your cluster
 - If a Pod crashes, it will be restarted
 - If a Node crashes, the Pods running on it will be lost, as well



- Each Pod is assigned to a specific node in your cluster
 - If a Pod crashes, it will be restarted
 - If a Node crashes, the Pods running on it will be lost, as well



High-Level Controllers



- ReplicaSet
 - Ensure that N instances of a given type of Pod are running at all times
 - If a Pod is lost (Node failure), the ReplicaSet will create a new one


```
~ ➔ kubectl get pods
```

NAME	READY	STATUS	RESTARTS	AGE
test-rs-42xcf	1/1	Running	0	1m
test-rs-lspwk	1/1	Running	0	1m
test-rs-swzg2	1/1	Running	0	1m

```
~ ➔ kubectl delete pod test-rs-42xcf
```

pod "test-rs-42xcf" deleted

→ kubectl get pods

NAME	READY	STATUS	RESTARTS	AGE
test-rs-fv87f	1/1	Running	0	12s
test-rs-lspwk	1/1	Running	0	9m
test-rs-swzg2	1/1	Running	0	9m

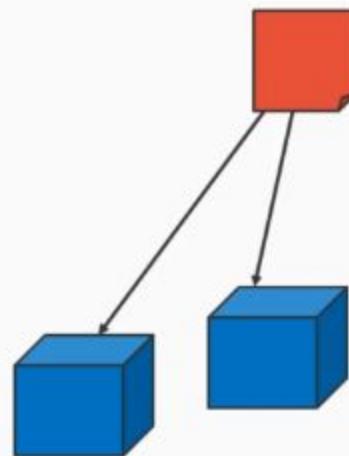
Conclusion



Single Pod

- disposable
- not resilient by itself
- not scalable by itself

Conclusion



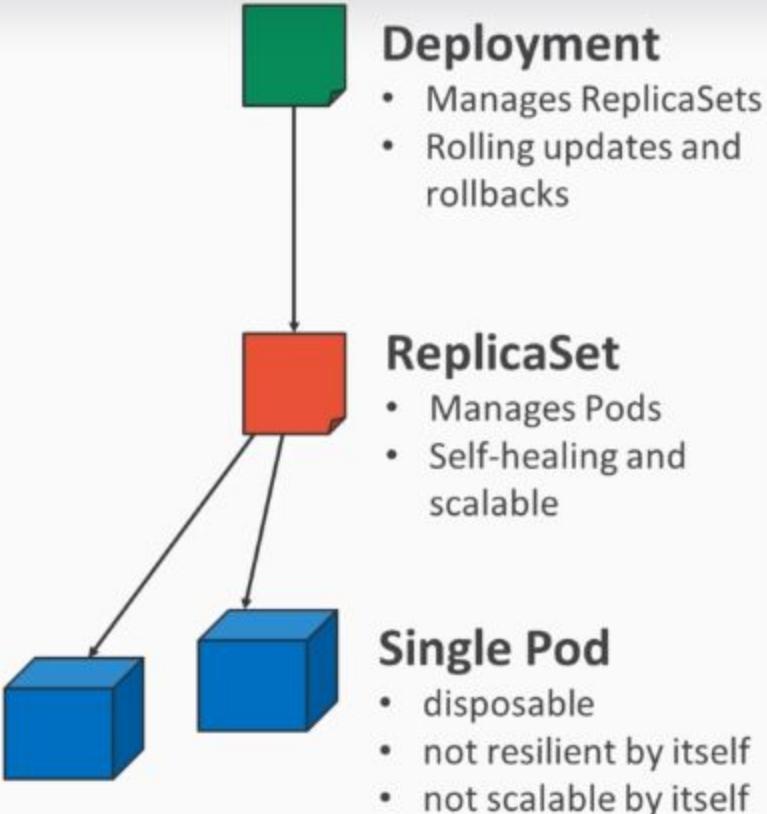
ReplicaSet

- Manages Pods
- Self-healing and scalable

Single Pod

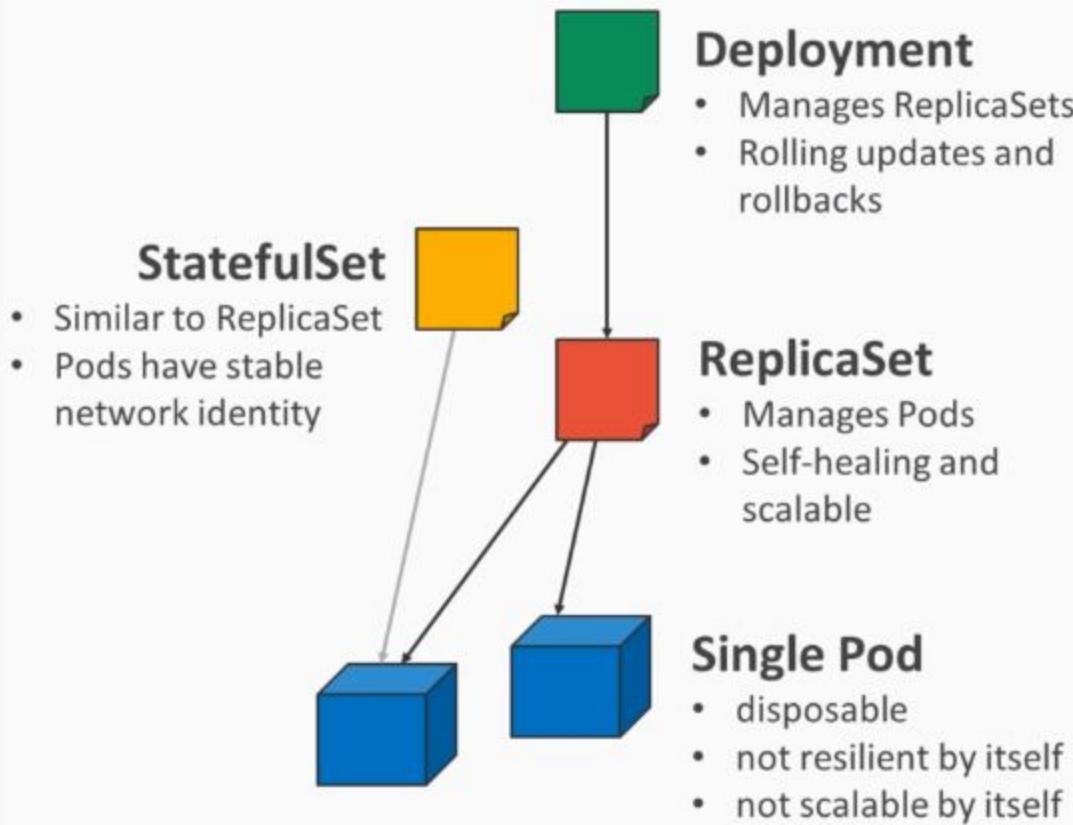
- disposable
- not resilient by itself
- not scalable by itself

Conclusion



Conclusion

- Rule of thumb: Never manually create a Pod
- Instead rely on controllers (ReplicaSet, Deployment, StatefulSet) to manage Pods for you
- More about ReplicaSets and Deployments in Section 2
- More about StatefulSets in Section 3



The Kubernetes Deployment Model

What Happened So Far?

```
mhelmich — mhelmich@Martins-MBP — ~ — -zsh
[~] ➜ kubectl scale replicaset test-rs --replicas=10
replicaset.extensions "test-rs" scaled
[~] ➜ kubectl get pods
NAME        READY   STATUS            RESTARTS   AGE
test-rs-5kwwr  1/1    Running          0          6s
test-rs-84mzk  0/1    ContainerCreating 0          6s
test-rs-bskk6  1/1    Running          0          45s
test-rs-ftpp5  1/1    Running          0          45s
test-rs-g8444  0/1    ContainerCreating 0          6s
test-rs-l4jt8  0/1    ContainerCreating 0          6s
test-rs-lshqd  0/1    ContainerCreating 0          6s
test-rs-vcn6b  0/1    ContainerCreating 0          6s
test-rs-xg9gb  1/1    Running          0          6s
test-rs-zfbkw  1/1    Running          0          45s
```

Labels

Label All the Things!

- In Kubernetes, every object can have arbitrary labels:

```
apiVersion: v1
kind: Pod
metadata:
  name: some-pod
  labels:
    app: my-great-microservice
    stage: production
```

Defining a ReplicaSet



ReplicaSet properties

Defining a ReplicaSet

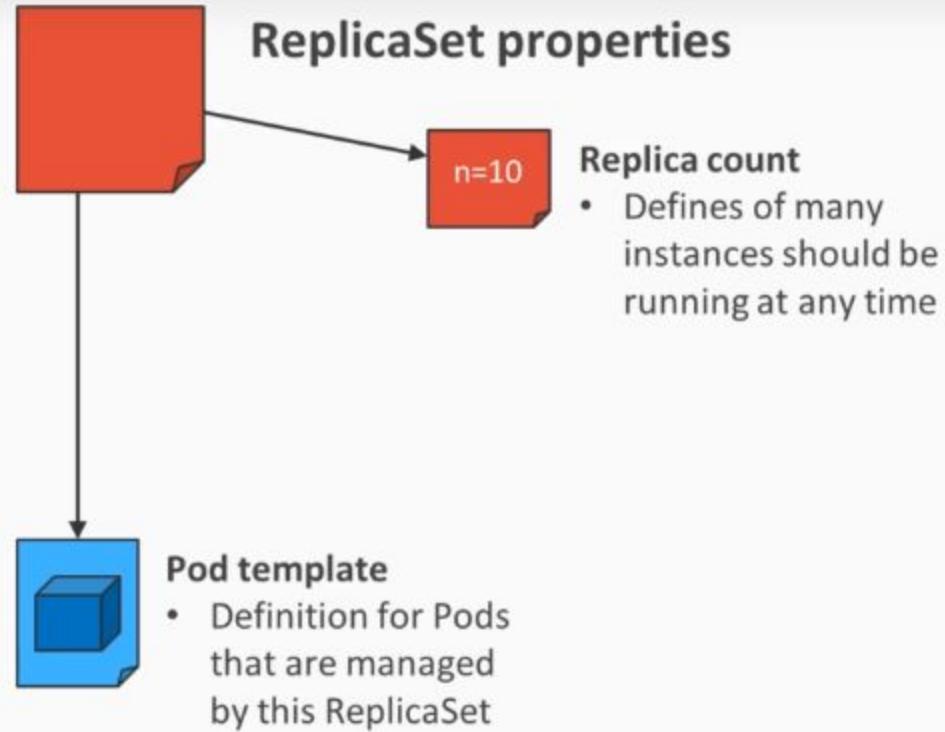


ReplicaSet properties

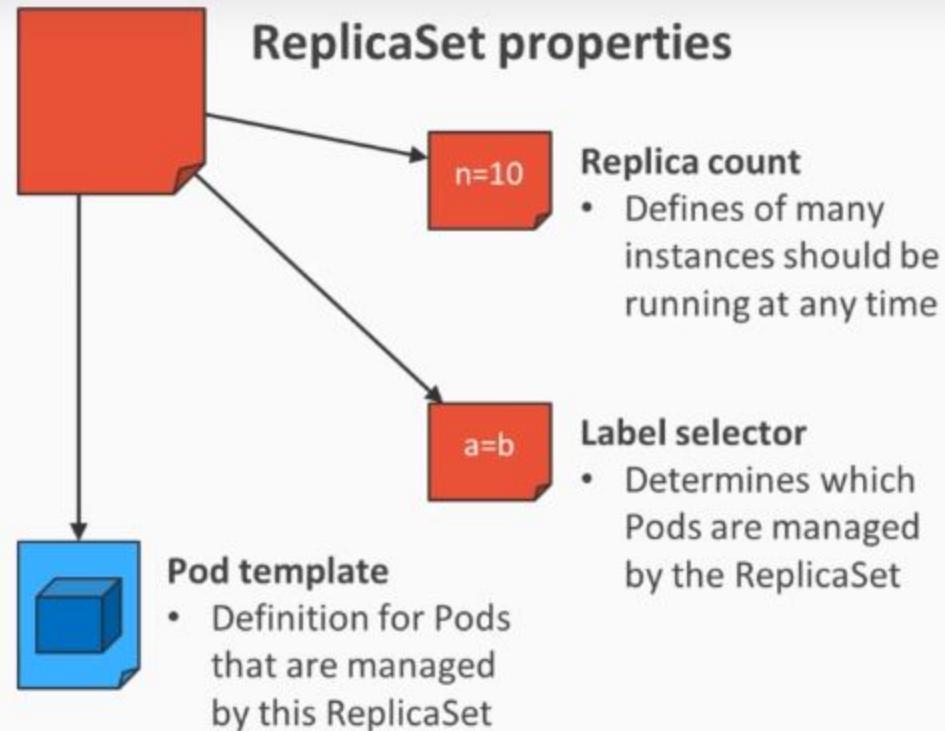
Pod template

- Definition for Pods that are managed by this ReplicaSet

Defining a ReplicaSet



Defining a ReplicaSet



ReplicaSet Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-service
  template:
    metadata:
      labels:
        app: my-service
    spec:
      containers:
        - name: service
          image: nginx
```

ReplicaSet Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-service ReplicaSet name
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-service
  template:
    metadata:
      labels:
        app: my-service
    spec:
      containers:
        - name: service
          image: nginx
```

ReplicaSet Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-service
spec:
  replicas: 3    Replica count
  selector:
    matchLabels:
      app: my-service
  template:
    metadata:
      labels:
        app: my-service
    spec:
      containers:
        - name: service
          image: nginx
```

ReplicaSet Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-service
template:
  metadata:
    labels:
      app: my-service
  spec:
    containers:
    - name: service
      image: nginx
```

Template used for
creating Pods
(same Pod specification as
in previous video)

ReplicaSet Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: my-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-service
  template:
    metadata:
      labels:
        app: my-service
  spec:
    containers:
      - name: service
        image: nginx
```

Label selector
Used by the ReplicaSet to determine which Pods are managed by this ReplicaSet

Pod metadata
MUST contain the selected labels from the selector

```
~/Kubernetes/2.1 Using ReplicaSets/Code ➤ kubectl apply -f replicaset.yaml
replicaset.apps "my-service" created
~/Kubernetes/2.1 Using ReplicaSets/Code ➤ kubectl get pods
NAME        READY   STATUS            RESTARTS   AGE
my-service-5bdmq  1/1    Running          0          14s
my-service-bw7cd  0/1    ContainerCreating  0          14s
my-service-tqvc8  0/1    ContainerCreating  0          14s
~/Kubernetes/2.1 Using ReplicaSets/Code ➤ kubectl scale replicaset my-service --replicas=10
replicaset.extensions "my-service" scaled
~/Kubernetes/2.1 Using ReplicaSets/Code ➤ kubectl get pods
NAME        READY   STATUS            RESTARTS   AGE
my-service-2rbpj  0/1    ContainerCreating  0          13s
my-service-5bdmq  1/1    Running          0          47s
my-service-bw7cd  1/1    Running          0          47s
my-service-fzvpx  1/1    Running          0          13s
my-service-n2cqh  1/1    Running          0          13s
my-service-ptbln  1/1    Running          0          13s
my-service-qxjkl  0/1    ContainerCreating  0          13s
my-service-ss9cg  0/1    ContainerCreating  0          13s
my-service-tqvc8  1/1    Running          0          47s
my-service-x2szj  1/1    Running          0          13s
~/Kubernetes/2.1 Using ReplicaSets/Code ➤
```

Connecting Services

What Happened So Far?

```
mhelmich — mhelmich@Martins-MBP — ~ — -zsh
[ ~ ➜ kubectl scale replicaset test-rs --replicas=10
replicaset.extensions "test-rs" scaled
[ ~ ➜ kubectl get pods
NAME        READY   STATUS    RESTARTS   AGE
test-rs-5kwwr  1/1     Running   0          6s
test-rs-84mzk  0/1     ContainerCreating   0          6s
test-rs-bskk6  1/1     Running   0          45s
test-rs-ftpp5  1/1     Running   0          45s
test-rs-g8444  0/1     ContainerCreating   0          6s
test-rs-l4jt8  0/1     ContainerCreating   0          6s
test-rs-lshqd  0/1     ContainerCreating   0          6s
test-rs-vcn6b  0/1     ContainerCreating   0          6s
test-rs-xg9gb  1/1     Running   0          6s
test-rs-zfbkw  1/1     Running   0          45s
```

Pod Networking



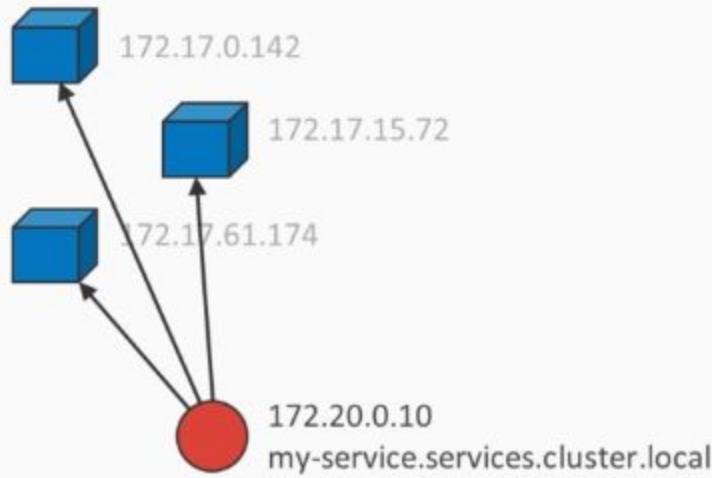
Pods

- Have IP addresses
- These are routed within the cluster (no external connectivity)

But

- Pods are short-lived
- Addresses may change frequently
- No DNS for Pod IPs
- No load balancing

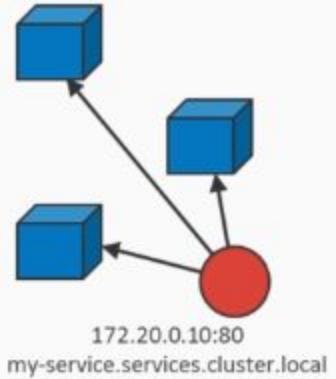
Services to the Rescue!



Services

- Selects a set of Pods by label
- One stable IP and DNS name for that group of Pods
- Round-robin load balancing
- Automatically adds/removes Pods (for example, managed by ReplicaSet)

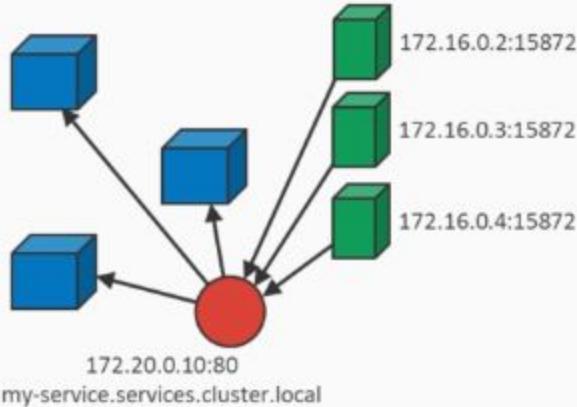
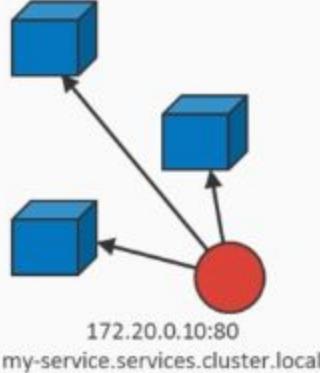
Service Types



ClusterIP

- Stable internal IP
- Stable internal DNS

Service Types



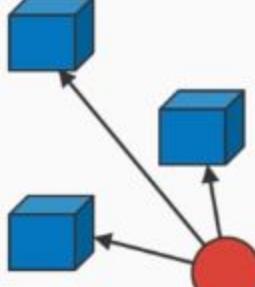
ClusterIP

- Stable internal IP
- Stable internal DNS

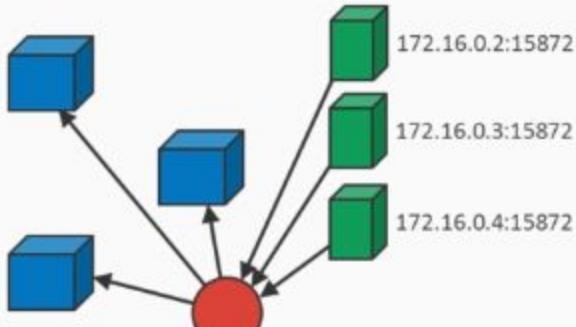
NodePort

- As ClusterIP, but additionally
- Every Node gets a public TCP port forwarding to that service

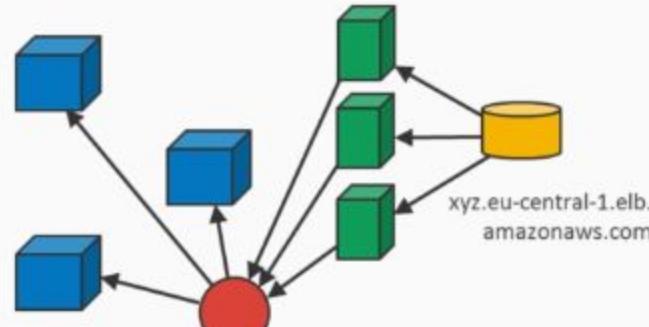
Service Types



172.20.0.10:80
my-service.services.cluster.local



172.20.0.10:80
my-service.services.cluster.local



172.20.0.10:80
my-service.services.cluster.local

ClusterIP

- Stable internal IP
- Stable internal DNS

NodePort

- As ClusterIP, but additionally
- Every Node gets a public TCP port forwarding to that service

LoadBalancer

- As NodePort, but additionally
- A load balancer (must be supported by cloud provider) is created to allow external incoming traffic

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
  Service name; will also be the
  DNS name
spec:
  type: ClusterIP
  selector:
    app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP Service type
  selector:
    app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

selector:
app: my-service

Label selector
Used to determine which Pods traffic will be forwarded to

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  type: ClusterIP
  selector:
    app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

Ports that should be forwarded by the service (supports TCP+UDP)

```
~/Kubernetes/2.2 Connecting to Services ➤ kubectl apply -f Code/service.yaml
service "my-service" created
~/Kubernetes/2.2 Connecting to Services ➤ kubectl get services
NAME      TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
kubernetes  ClusterIP  10.96.0.1    <none>        443/TCP      19m
my-service  ClusterIP  10.110.126.238  <none>        80/TCP      15s
~/Kubernetes/2.2 Connecting to Services ➤ kubectl run --rm -it --image=alpine my-test
If you don't see a command prompt, try pressing enter.
/ # apk -U add curl
fetch http://dl-cdn.alpinelinux.org/alpine/v3.7/main/x86_64/APKINDEX.tar.gz
fetch http://dl-cdn.alpinelinux.org/alpine/v3.7/community/x86_64/APKINDEX.tar.gz
(1/4) Installing ca-certificates (20171114-r0)
(2/4) Installing libssh2 (1.8.0-r2)
(3/4) Installing libcurl (7.60.0-r1)
(4/4) Installing curl (7.60.0-r1)
Executing busybox-1.27.2-r7.trigger
Executing ca-certificates-20171114-r0.trigger
OK: 6 MiB in 15 packages
/ # curl -v my-service
```

```
< Content-Length: 612
< Last-Modified: Mon, 09 Apr 2018 16:01:09 GMT
< Connection: keep-alive
< ETag: "5acb8e45-264"
< Accept-Ranges: bytes
<
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
body {
    width: 35em;
    margin: 0 auto;
    font-family: Tahoma, Verdana, Arial, sans-serif;
}
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

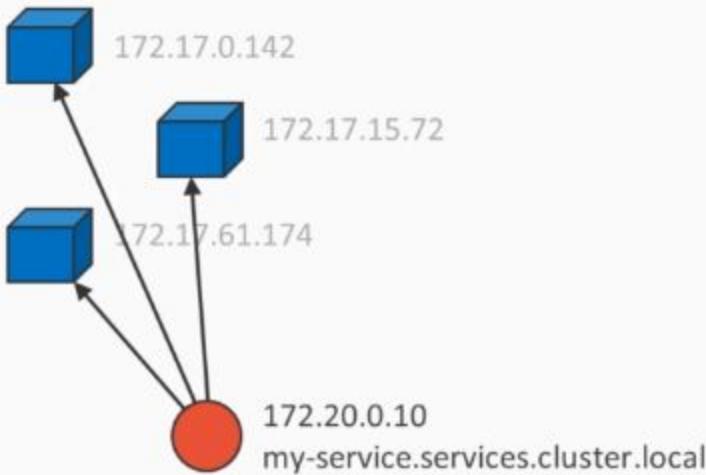
<p><em>Thank you for using nginx.</em></p>
</body>
</html>
* Connection #0 to host my-service left intact
/ # █
```

Service Example

```
apiVersion: v1
kind: Service
metadata:
  name: my-service
spec:
  - type: ClusterIP
  + type: LoadBalancer
    selector:
      app: my-service
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
```

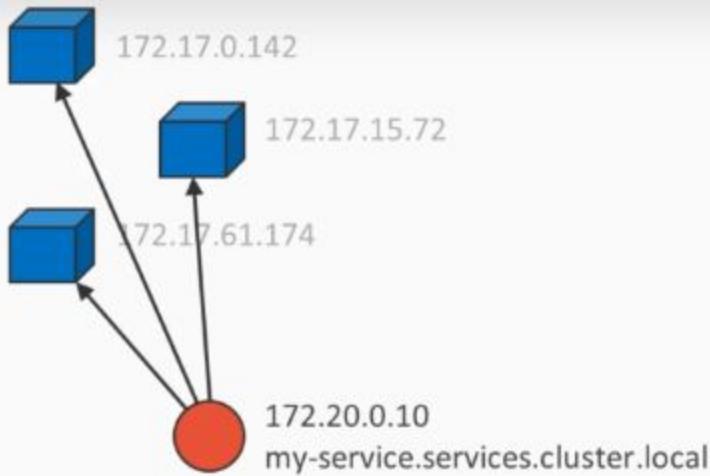
HTTP Connectivity with Ingress

Previously on



Services

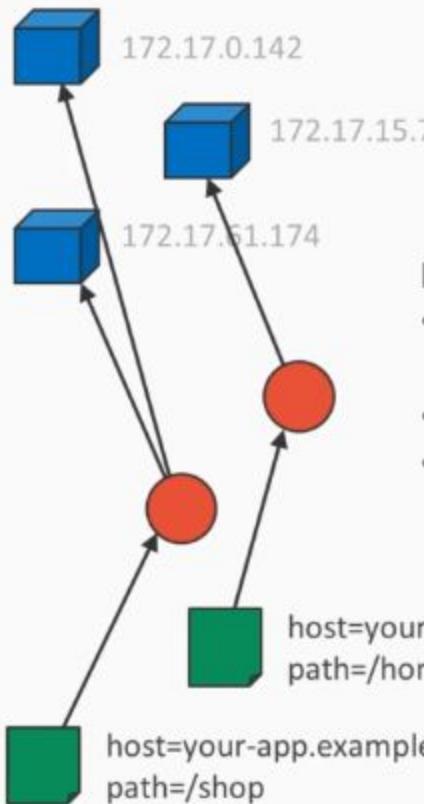
Service Limitations



Services

- Do raw TCP/UDP forwarding
- Do not know about application protocols like HTTP

Ingress Resources



Ingress

- Defines request routing for HTTP requests to services
- Works at application layer
- Supports path and host-based routing, caching, authentication, and more

Ingress Example

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-ingress
spec:
  rules:
    - host: your-app.example
      http:
        paths:
          - path: /shop
            backend:
              serviceName: my-service
              servicePort: 80
          - path: /home
            backend:
              serviceName: my-other-
service
              servicePort: 80
```

Ingress Example

```
apiVersion: extensions/v1beta1
kind: Ingress
metadata:
  name: my-ingress
spec:                                     Host rules; one per HTTP host name
  rules:
    - host: your-app.example
      http:
        paths:
          - path: /shop
            backend:
              serviceName: my-service
              servicePort: 80
          - path: /home
            backend:
              serviceName: my-other-
              service
              servicePort: 80
```

Ingress Controllers



host=your-app.example
path=/home

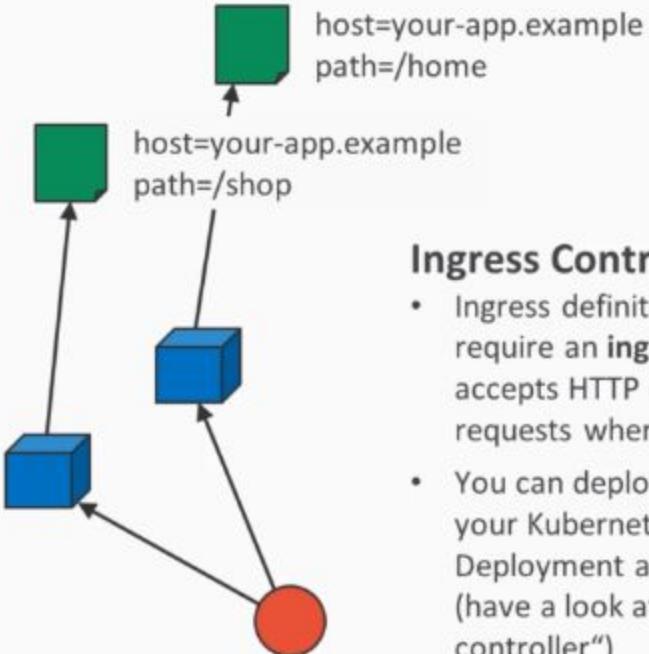


host=your-app.example
path=/shop

Ingress Controllers

- Ingress definitions are virtual; they require an **ingress controller** that accepts HTTP connections and forwards requests where they should go

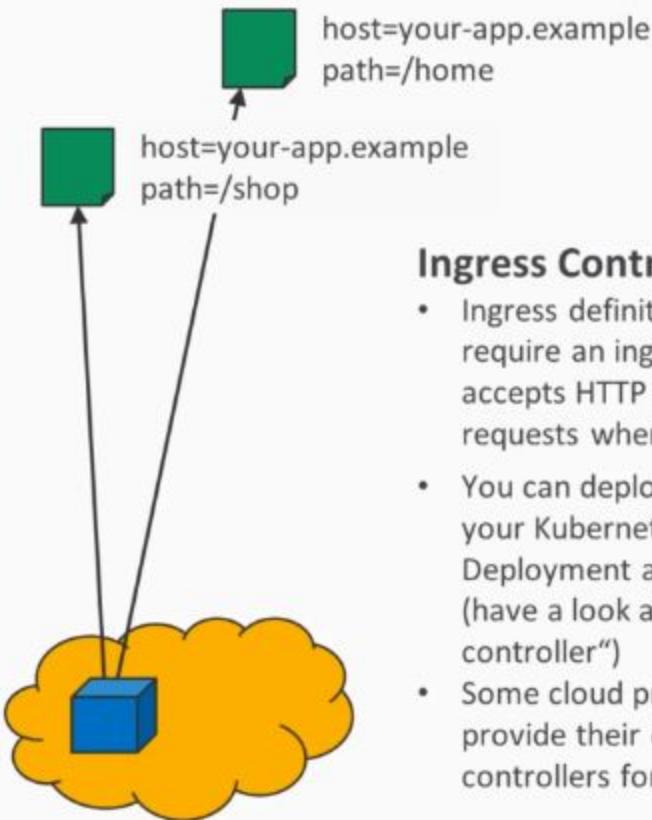
Ingress Controllers



Ingress Controllers

- Ingress definitions are virtual; they require an **ingress controller** that accepts HTTP connections and forwards requests where they should go
- You can deploy an Ingress controller in your Kubernetes cluster, using regular Deployment and Service definitions (have a look at the “NGINX ingress controller”)

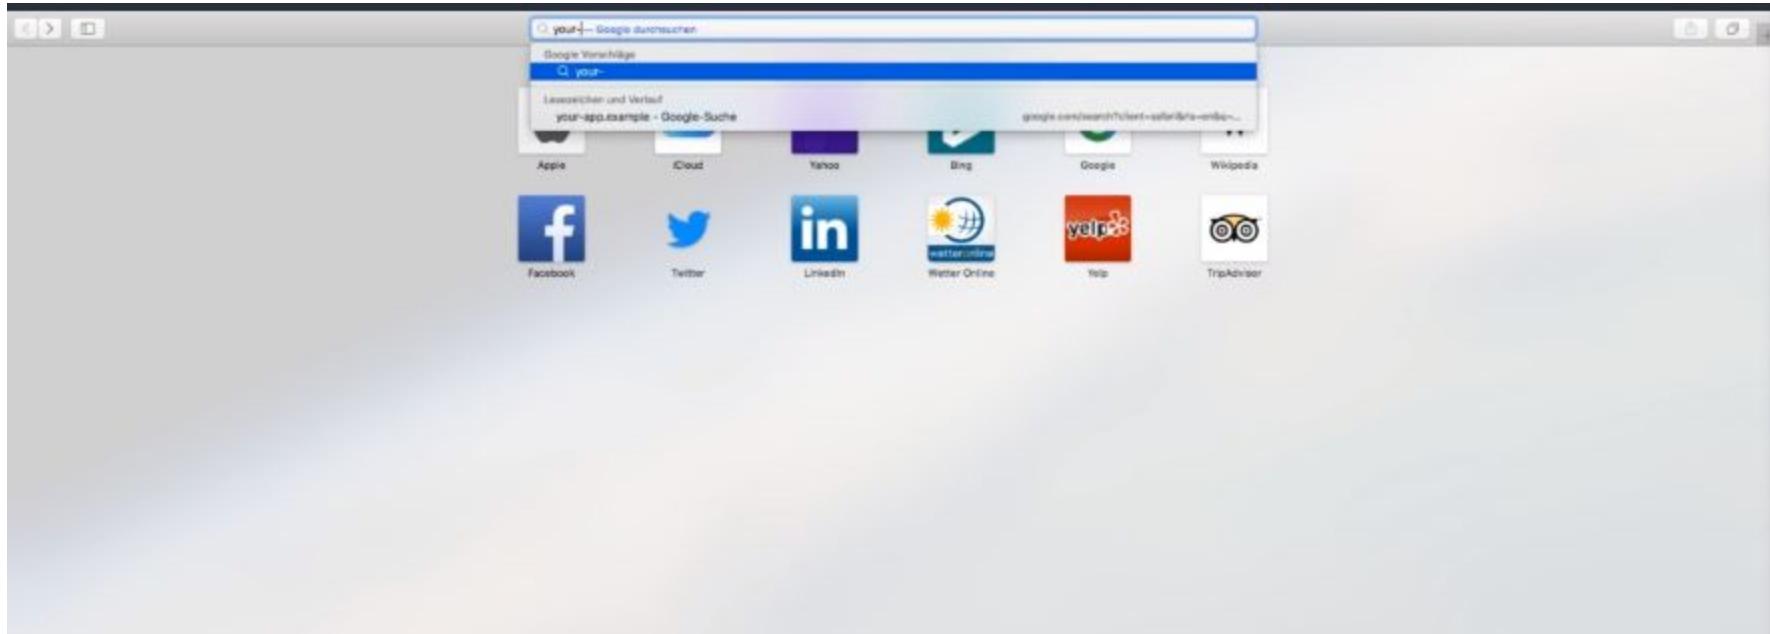
Ingress Controllers



Ingress Controllers

- Ingress definitions are virtual; they require an ingress controller that accepts HTTP connections and forwards requests where they should go
- You can deploy an Ingress controller in your Kubernetes clusters, using regular Deployment and Service definitions (have a look at the “NGINX ingress controller”)
- Some cloud providers (like GKE) also provide their own ready-to-use Ingress controllers for you

```
~/Kubernetes/2.3 Using Ingress Resources/Code ➤ kubectl apply -f ingress.yaml
ingress.extensions "my-ingress" unchanged
~/Kubernetes/2.3 Using Ingress Resources/Code ➤ minikube addons list
- addon-manager: enabled
- coredns: disabled
- dashboard: disabled
- default-storageclass: enabled
- efk: disabled
- freshpod: disabled
- heapster: enabled
- ingress: enabled
- kube-dns: enabled
- registry: disabled
- registry-creds: disabled
- storage-provisioner: enabled
~/Kubernetes/2.3 Using Ingress Resources/Code ➤ minikube addons enable ingress
Ingress was successfully enabled
~/Kubernetes/2.3 Using Ingress Resources/Code ➤ minikube ip
192.168.99.100
~/Kubernetes/2.3 Using Ingress Resources/Code ➤ sudo vim /etc/hosts
Password:
```



Ingress Controller Features

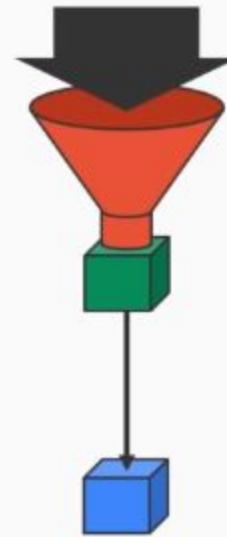


TLS Offloading

Ingress Controller Features



TLS Offloading

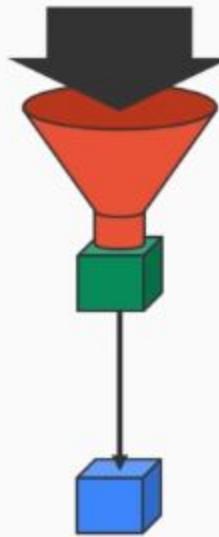


Rate Limiting

Ingress Controller Features



TLS Offloading



Rate Limiting



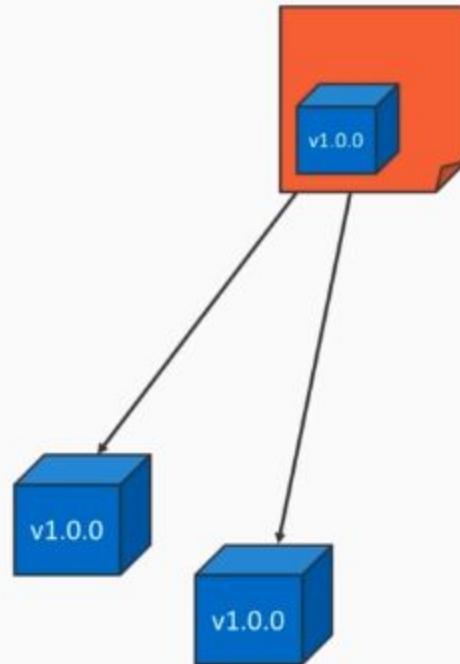
Authentication

Managing the Application Lifecycle with Deployments

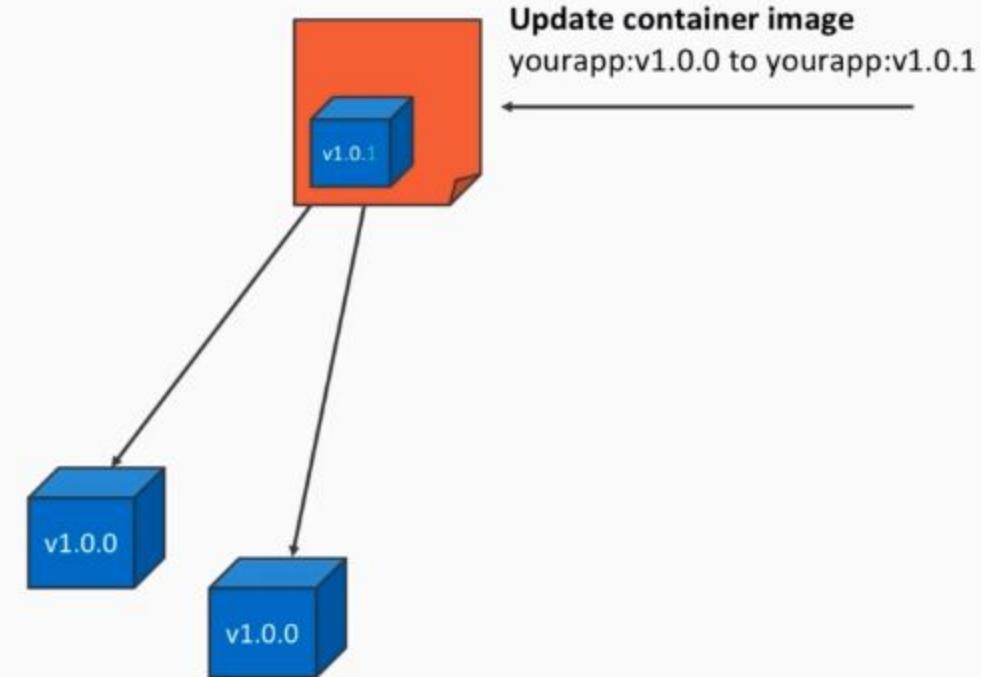
Previously on

```
mhelmich — mhelmich@Martins-MBP — ~ — -zsh
[~] ➜ kubectl scale replicaset test-rs --replicas=10
replicaset.extensions "test-rs" scaled
[~] ➜ kubectl get pods
NAME      READY   STATUS        RESTARTS   AGE
test-rs-5kwwr  1/1    Running     0          6s
test-rs-84mzk   0/1    ContainerCreating 0          6s
test-rs-bskk6   1/1    Running     0          45s
test-rs-ftpp5   1/1    Running     0          45s
test-rs-g8444   0/1    ContainerCreating 0          6s
test-rs-l4jt8   0/1    ContainerCreating 0          6s
test-rs-lshqd   0/1    ContainerCreating 0          6s
test-rs-vcn6b   0/1    ContainerCreating 0          6s
test-rs-xg9gb   1/1    Running     0          6s
test-rs-zfbkw   1/1    Running     0          45s
[~] ➜
```

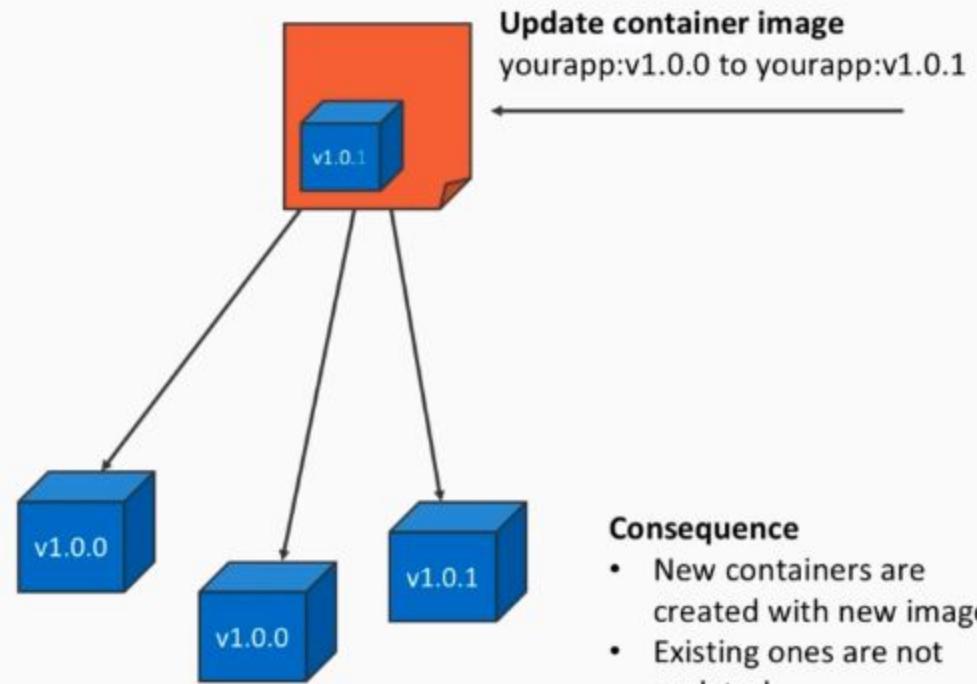
ReplicaSet Limitations



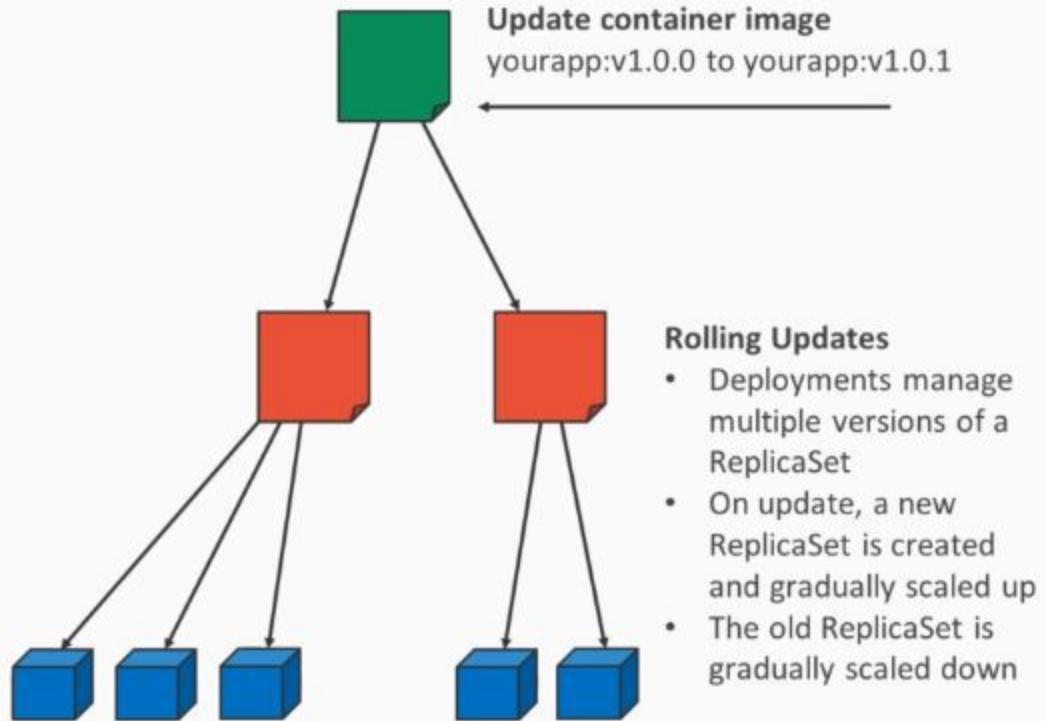
ReplicaSet Limitations



ReplicaSet Limitations



Solution: Deployments



Deployment Example

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: my-service
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-service
  template:
    metadata:
      labels:
        app: my-service
    spec:
      containers:
        - name: service
          image: nginx:1.13.11
```

Hey! This looks
just like a
ReplicaSet!

```
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl apply -f deployment.yaml
deployment.apps "my-deployment" created
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get pods
NAME                               READY   STATUS            RESTARTS   AGE
my-deployment-747ddbf996-hkds4    0/1    ContainerCreating   0          5s
my-deployment-747ddbf996-jvlpd    0/1    ContainerCreating   0          5s
my-deployment-747ddbf996-nrtvp    0/1    ContainerCreating   0          5s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get replicsets
NAME           DESIRED   CURRENT   READY   AGE
my-deployment-747ddbf996   3         3        3       26s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl edit deployment my-deployment
```

```
creationTimestamp: 2018-05-31T11:22:14Z
generation: 1
labels:
  app: my-service
name: my-deployment
namespace: default
resourceVersion: "9864"
selfLink: /apis/extensions/v1beta1/namespaces/default/deployments/my-deployment
uid: de6cb09b-64c4-11e8-8174-080027c553e0
spec:
  progressDeadlineSeconds: 600
  replicas: 3
  revisionHistoryLimit: 10
  selector:
    matchLabels:
      app: my-service
  strategy:
    rollingUpdate:
      maxSurge: 25%
      maxUnavailable: 25%
    type: RollingUpdate
  template:
    metadata:
      creationTimestamp: null
      labels:
        app: my-service
    spec:
      containers:
        - image: nginx:1.13.12
          imagePullPolicy: IfNotPresent
          name: nginx
          resources: {}
```

```
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl apply -f deployment.yaml
deployment.apps "my-deployment" created
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddbf996-hkds4  0/1    ContainerCreating   0          5s
my-deployment-747ddbf996-jvlpd  0/1    ContainerCreating   0          5s
my-deployment-747ddbf996-nrtvp  0/1    ContainerCreating   0          5s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get replicaset
NAME           DESIRED   CURRENT   READY   AGE
my-deployment-747ddbf996   3         3         3      26s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl edit deployment my-deployment
deployment.extensions "my-deployment" edited
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddbf996-hkds4  1/1    Running     0          1m
my-deployment-747ddbf996-jvlpd  1/1    Running     0          1m
my-deployment-747ddbf996-nrtvp  1/1    Running     0          1m
my-deployment-7dd8f6c5b5-8pnvv  0/1    ContainerCreating   0          2s
~/Kubernetes/2.4 Using Deployments/Code ➤
```

```
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl apply -f deployment.yaml
deployment.apps "my-deployment" created
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddb996-hkds4  0/1    ContainerCreating   0          5s
my-deployment-747ddb996-jvlpd  0/1    ContainerCreating   0          5s
my-deployment-747ddb996-nrtvp  0/1    ContainerCreating   0          5s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
my-deployment-747ddb996   3         3         3       26s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl edit deployment my-deployment
deployment.extensions "my-deployment" edited
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddb996-hkds4  1/1    Running     0          1m
my-deployment-747ddb996-jvlpd  1/1    Running     0          1m
my-deployment-747ddb996-nrtvp  1/1    Running     0          1m
my-deployment-7dd8f6c5b5-8pnvv  0/1    ContainerCreating   0          2s
~/Kubernetes/2.4 Using Deployments/Code ➤ kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
my-deployment-747ddb996   0         0         0       1m
my-deployment-7dd8f6c5b5   3         3         3       21s
~/Kubernetes/2.4 Using Deployments/Code ➤
```

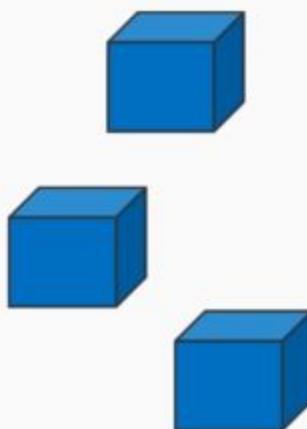
```
deployment.apps "my-deployment" created
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddbf996-hkds4  0/1    ContainerCreating  0          5s
my-deployment-747ddbf996-jvlpd  0/1    ContainerCreating  0          5s
my-deployment-747ddbf996-nrtvp  0/1    ContainerCreating  0          5s
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
my-deployment-747ddbf996   3         3         3       26s
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl edit deployment my-deployment
deployment.extensions "my-deployment" edited
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddbf996-hkds4  1/1    Running     0          1m
my-deployment-747ddbf996-jvlpd  1/1    Running     0          1m
my-deployment-747ddbf996-nrtvp  1/1    Running     0          1m
my-deployment-7dd8f6c5b5-8pnvv  0/1    ContainerCreating  0          2s
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
my-deployment-747ddbf996   0         0         0       1m
my-deployment-7dd8f6c5b5   3         3         3       21s
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl rollout undo deployment/my-deployment
deployment.apps "my-deployment"
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get replicaset
NAME      DESIRED   CURRENT   READY   AGE
my-deployment-747ddbf996   3         3         3       2m
my-deployment-7dd8f6c5b5   0         0         0       1m
~/Kubernetes/2.4 Using Deployments/Code ➔ kubectl get pods
NAME                      READY   STATUS        RESTARTS   AGE
my-deployment-747ddbf996-6hdpn  1/1    Running     0          15s
my-deployment-747ddbf996-8jjlh  1/1    Running     0          16s
my-deployment-747ddbf996-m8prz  1/1    Running     0          18s
~/Kubernetes/2.4 Using Deployments/Code ➔
```

Stateful Services and Storage Drivers

Up Until Now

Stateless Pods

- Pods that do not have any kind of state that needs to be stored
- Data is stored/retrieved from external storages

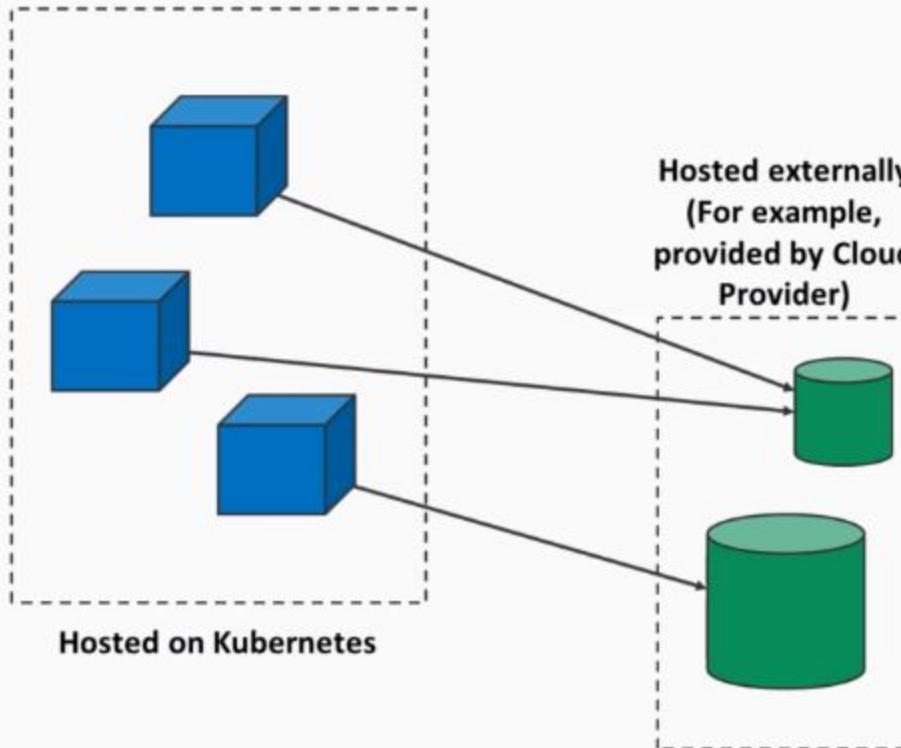


Hosted on Kubernetes

Up Until Now

Stateless Pods

- Pods that do not have any kind of state that needs to be stored
- Data is stored/retrieved from external storages



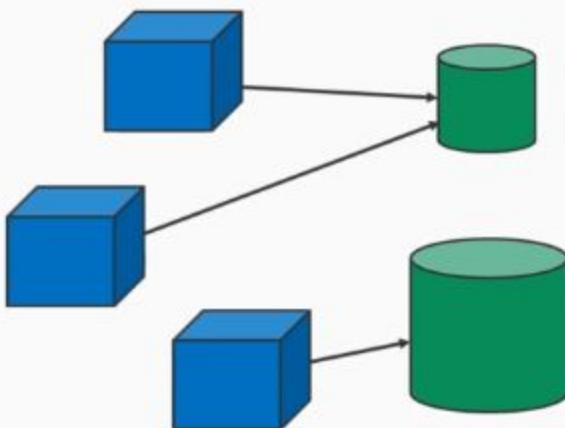
External Storage Services

- Databases
- Object Stores (like AWS S3 and similar)

Upcoming

Stateless Pods

- Pods that do not have any kind of state that needs to be stored
- Data is stored/retrieved from external storages

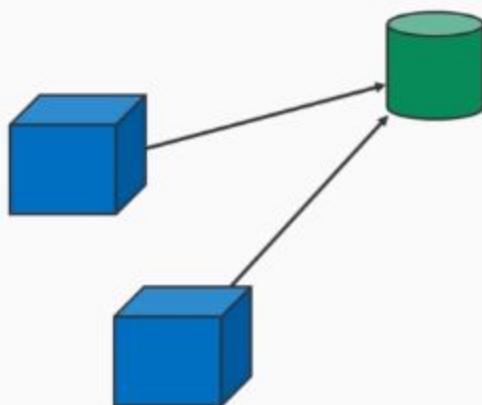


Stateful Pods

- Pods that have state/need to store data
- Require persistent file system across Pod termination/recreation

Hosted on Kubernetes

Persistent Volumes (Short: PV)



Persistent Volume

- Can be provided as a file system to Kubernetes Pods
- Lifecycle is independent from Pods
- Can be accessed by multiple Pods at once (dependent on storage engine)



Choosing a Network Storage Technology

Network Storage Technologies

- Network Block Devices

- Examples (cloud providers):

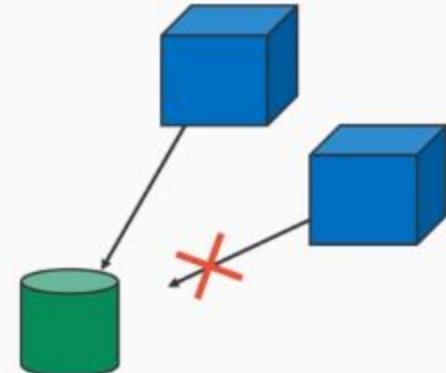
- AWS Elastic Block Devices
 - GCE Persistent Disk
 - AzureDisk
 - OpenStack Cinder

- Examples (self-hosting):

- Rados Block Devices (Ceph)
 - iSCSI

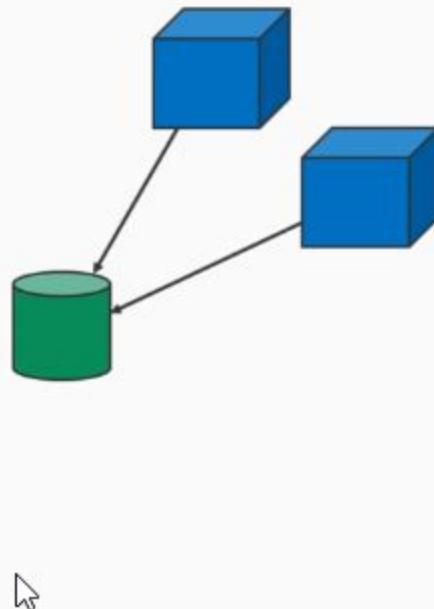
- Key features

- Fast
 - Fixed-size
 - Can (usually) be used by one Pod at a time



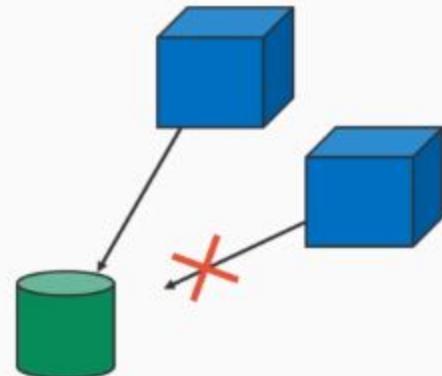
Network Storage Technologies (Continued)

- Network File Systems
 - Examples (cloud providers):
 - AWS Elastic File System (actually NFS)
 - Azure Files (actually SMB)
 - Examples (self-hosting):
 - CephFS
 - GlusterFS
 - NFS
 - Key features
 - Slower than Block Devices
 - Can be used by many Pods at a time



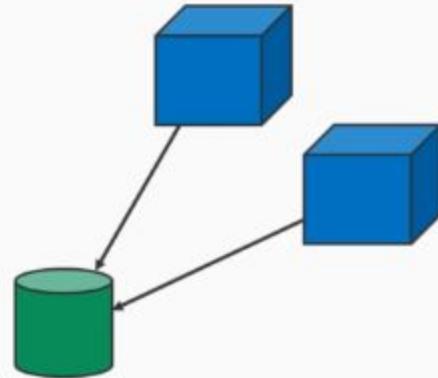
Network Storage Technologies

- Network Block Devices
 - Examples (cloud providers):
 - AWS Elastic Block Devices
 - GCE Persistent Disk
 - AzureDisk
 - OpenStack Cinder
 - Examples (self-hosting):
 - Rados Block Devices (Ceph)
 - iSCSI
 - Key features
 - Fast
 - Fixed-size
 - Can (usually) be used by one Pod at a time



Network Storage Technologies (Continued)

- Network File Systems
 - Examples (cloud providers):
 - AWS Elastic File System (actually NFS)
 - Azure Files (actually SMB)
 - Examples (self-hosting):
 - CephFS
 - GlusterFS
 - NFS
 - Key features
 - Slower than Block Devices
 - Can be used by many Pods at a time



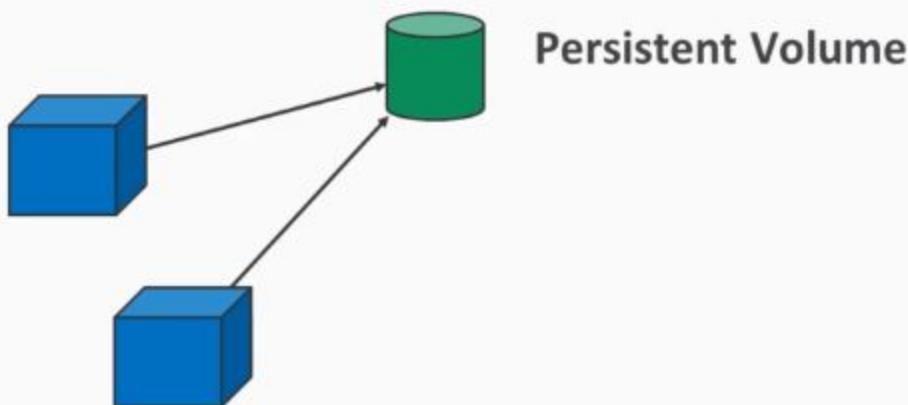
Final Notes

Storage Is Easy Hard

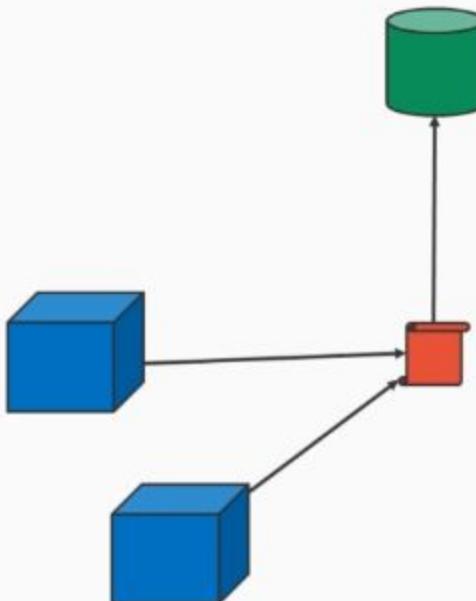
- Setting up Ceph or other self-hosted solutions is out-of-scope for this course
- Have a look at Learning Ceph by Anthony D'Atri, Vaibhav Bhembre and Karan Singh (<https://www.packtpub.com/virtualization-and-cloud/learning-ceph-second-edition>)
- Mixed operation between storage engines is possible

Working with Persistent Volumes

What Happened Previously?



Volumes and Claims



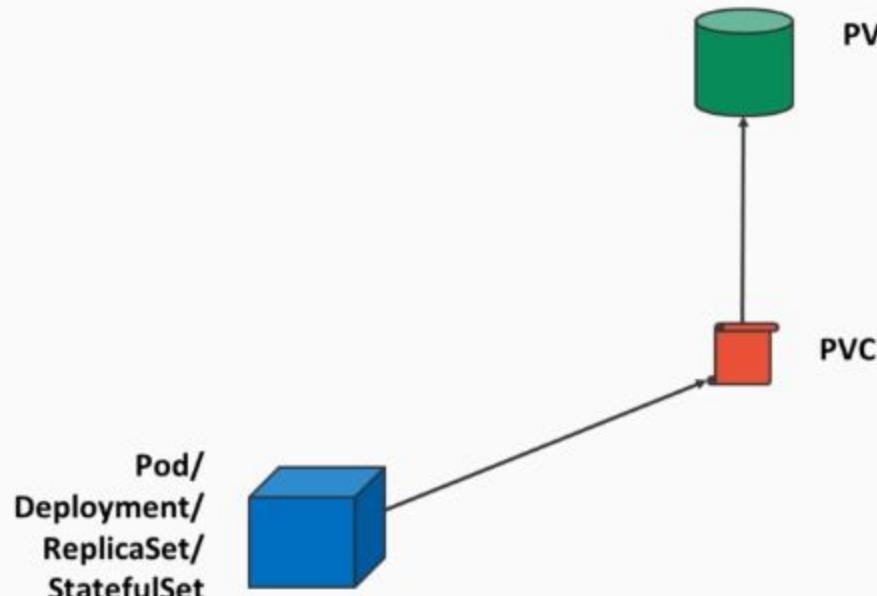
PersistentVolume

- Refers to a specific network block device or file system

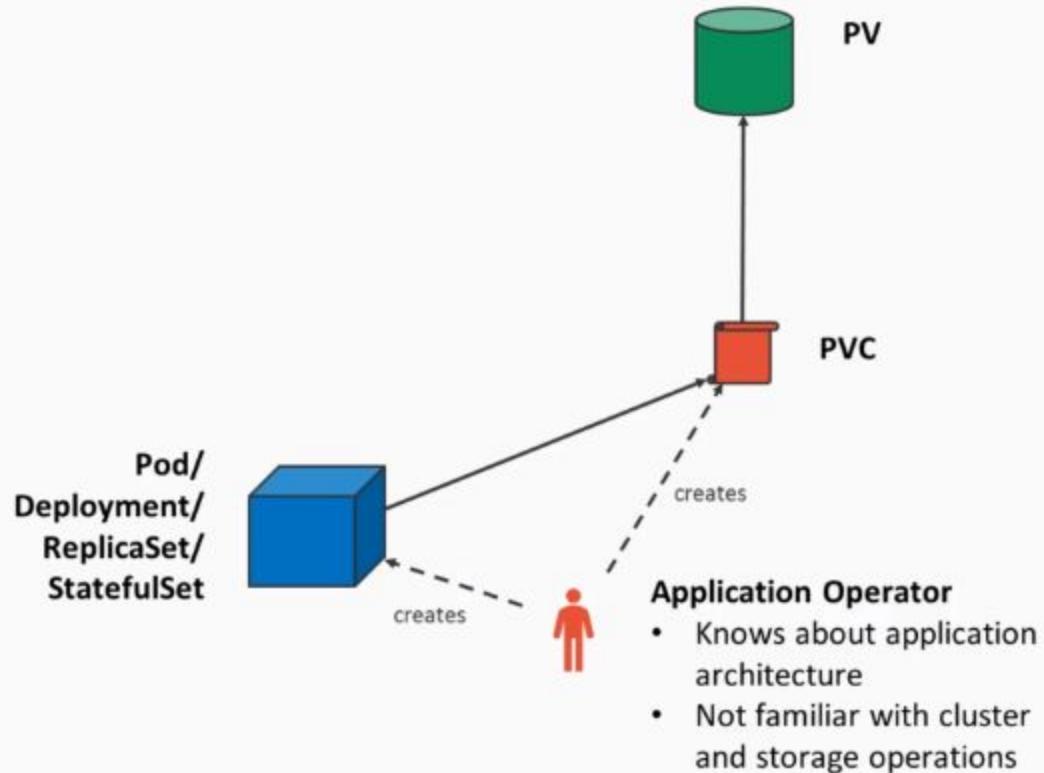
PersistentVolumeClaim

- Claims a PV for use in one or more Pods

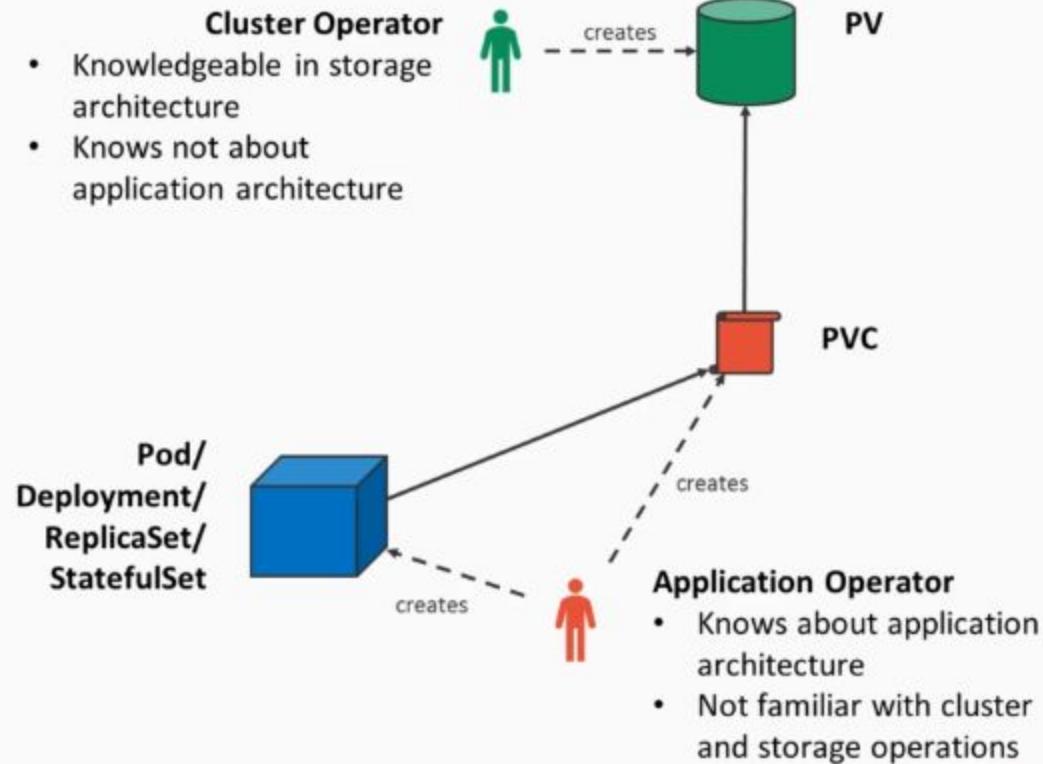
Why? Separation of Concerns!



Why? Separation of Concerns!



Why? Separation of Concerns!



PersistentVolume Example

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-volume
spec:
  capacity: Volume capacity
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  awsElasticBlockStore:
    volumeID: vol-08374e22abe6cb9b6
    fsType: ext4
```

PersistentVolume Example

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-volume
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  awsElasticBlockStore:
    volumeID: vol-08374e22abe6cb9b6
    fsType: ext4
```

Access modes:

- ReadWriteOnce – One Pod at once
- ReadWriteMany – Many Pods at once

PersistentVolume Example

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-volume
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  awsElasticBlockStore:
    volumeID: vol-08374e22abe6cb9b6
    fsType: ext4
```

Actual volume location; depends on
volume driver (AWS in this case)

PersistentVolume Example (For Minikube Users)

```
apiVersion: v1
kind: PersistentVolume
metadata:
  name: my-volume
spec:
  capacity:
    storage: 10Gi
  accessModes:
    - ReadWriteOnce
  hostPath:
    path: /data/my-volume
```

Caution!
Never use a
hostPath volume in
a clustered
environment!

PersistentVolumeClaim Example

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-volume-claim
spec:
  resources:
    requests:
      storage: 10Gi
  accessModes:
    - ReadWriteOnce
```

Required capacity

PersistentVolumeClaim Example

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-volume-claim
spec:
  resources:
    requests:
      storage: 10Gi
  accessModes:
    - ReadWriteOnce
Required access mode
```

Using Volumes in Pods

```
apiVersion: v1
kind: Pod
metadata:
  name: my-database
spec:
  containers:
    - name: database
      image: mysql:5.7
      env:
        # ...
+     volumeMounts:
+       - mountPath: /var/lib/mysql
+         name: data
+     volumes:
+       - name: data
+         persistentVolumeClaim:
+           claimName: my-volume-claim
```

Volume claims tha
a Pod should use

Using Volumes in Pods

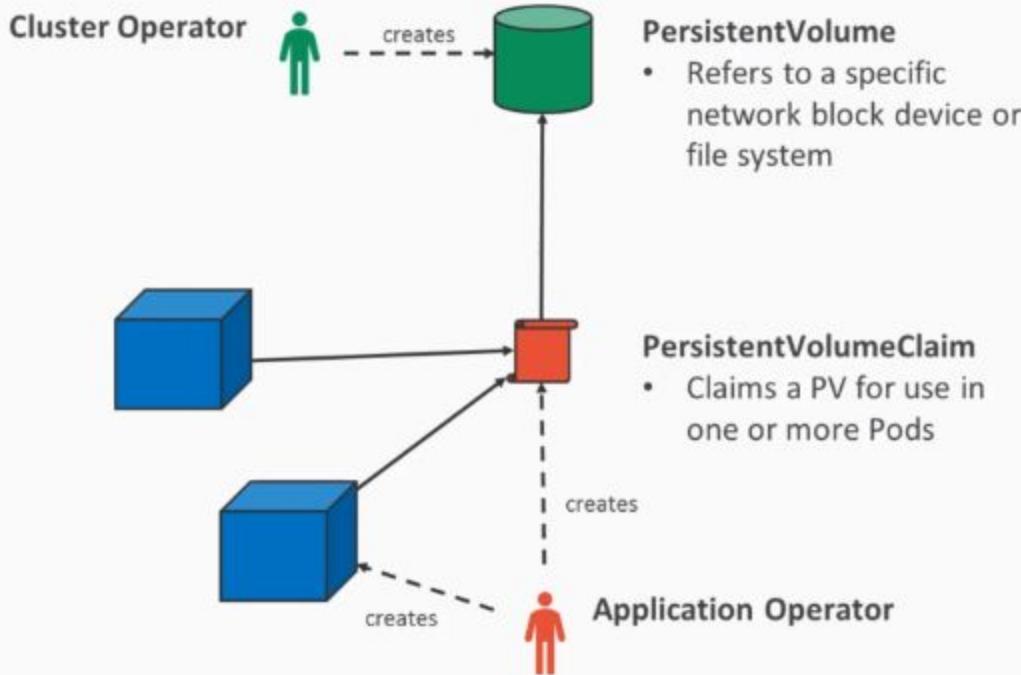
```
apiVersion: v1
kind: Pod
metadata:
  name: my-database
spec:
  containers:
    - name: database
      image: mysql:5.7
      env:
        # ...
      + volumeMounts:
        + - mountPath: /var/lib/mysql
          name: data
      + volumes:
        + - name: data
        + persistentVolumeClaim:
          claimName: my-volume-claim
```

Where to mount the volume within the Pod

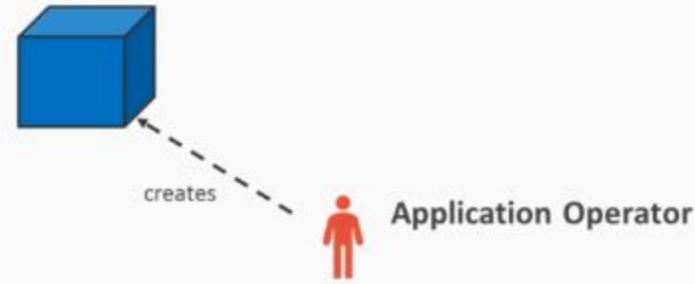
```
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl apply -f persistentvolume.yaml
persistentvolume "my-volume" created
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl get pv
NAME      CAPACITY   ACCESS MODES   RECLAIM POLICY   STATUS    CLAIM        STORAGECLASS   REASON   AGE
my-volume  10Gi       RWO          Retain          Available           22s
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl apply -f persistentvolumeclaim.yaml
persistentvolumeclaim "my-volume-claim" created
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl get pvc
NAME      STATUS    VOLUME          CAPACITY   ACCESS MODES   STORAGECLASS   AGE
my-volume-claim  Bound    pvc-7f9c4429-6777-11e8-a87a-080027c553e0  5Gi        RWO          standard      13s
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl apply -f pod.yaml
pod "my-database" created
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
my-database  1/1     Running   0          5s
~/Kubernetes/3.2 Working with Persistent Volumes/Code ➔
```

Automatic Volume Provisioning

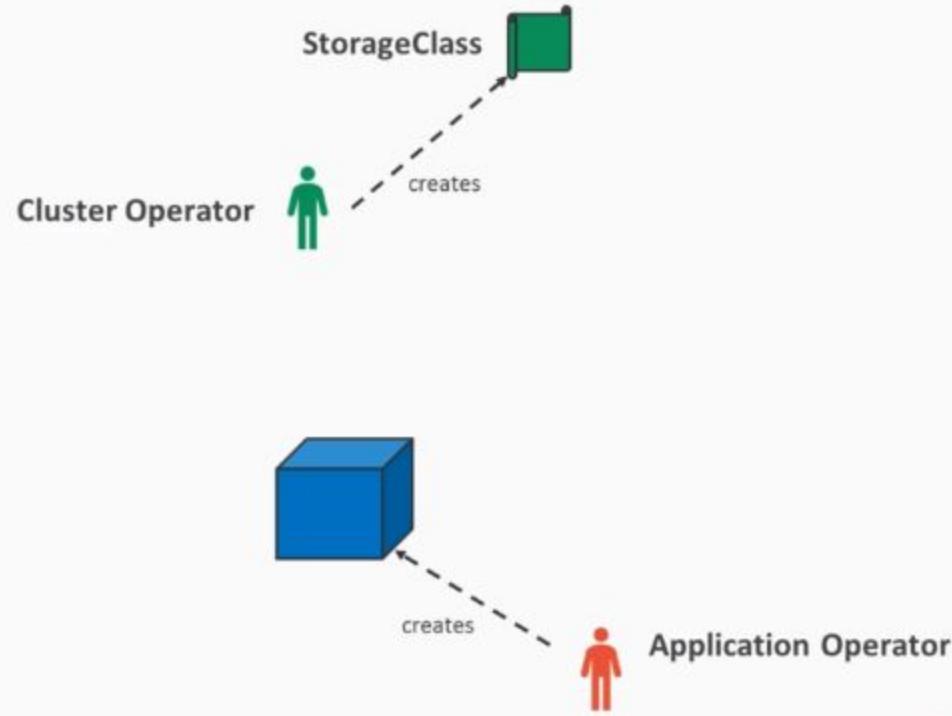
What Happened Previously?



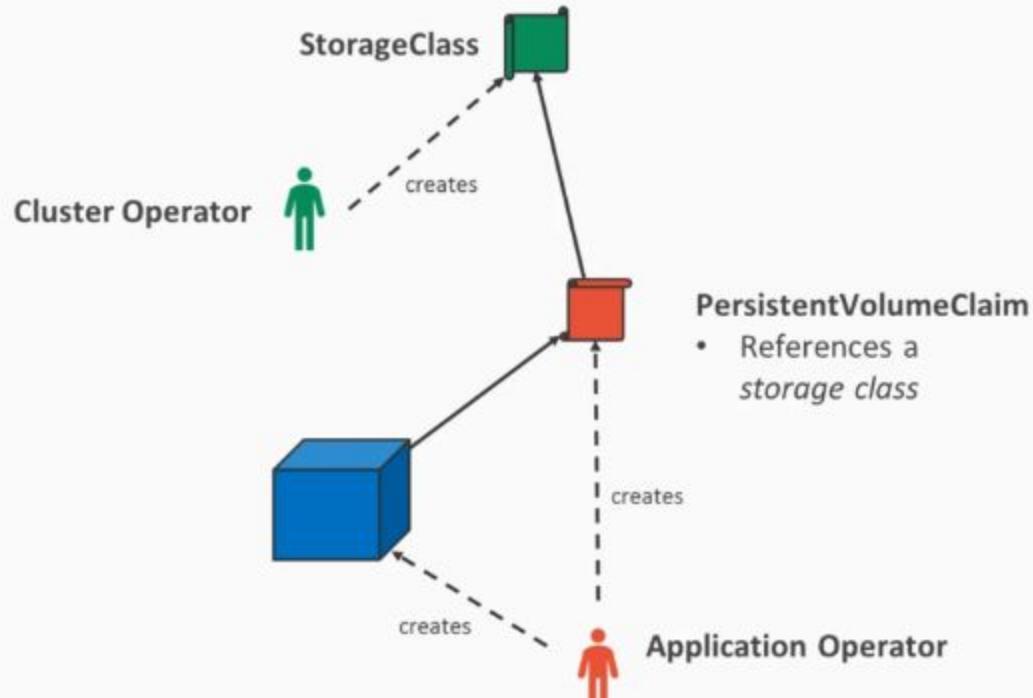
Storage Classes



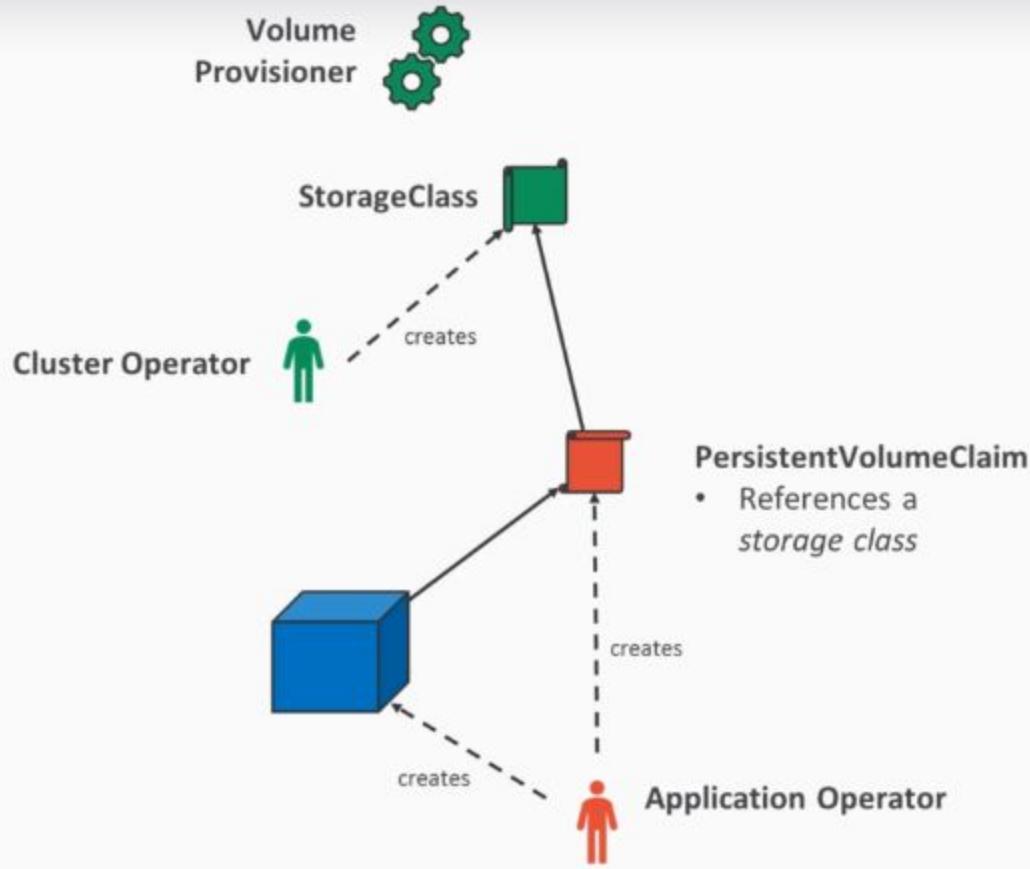
Storage Classes



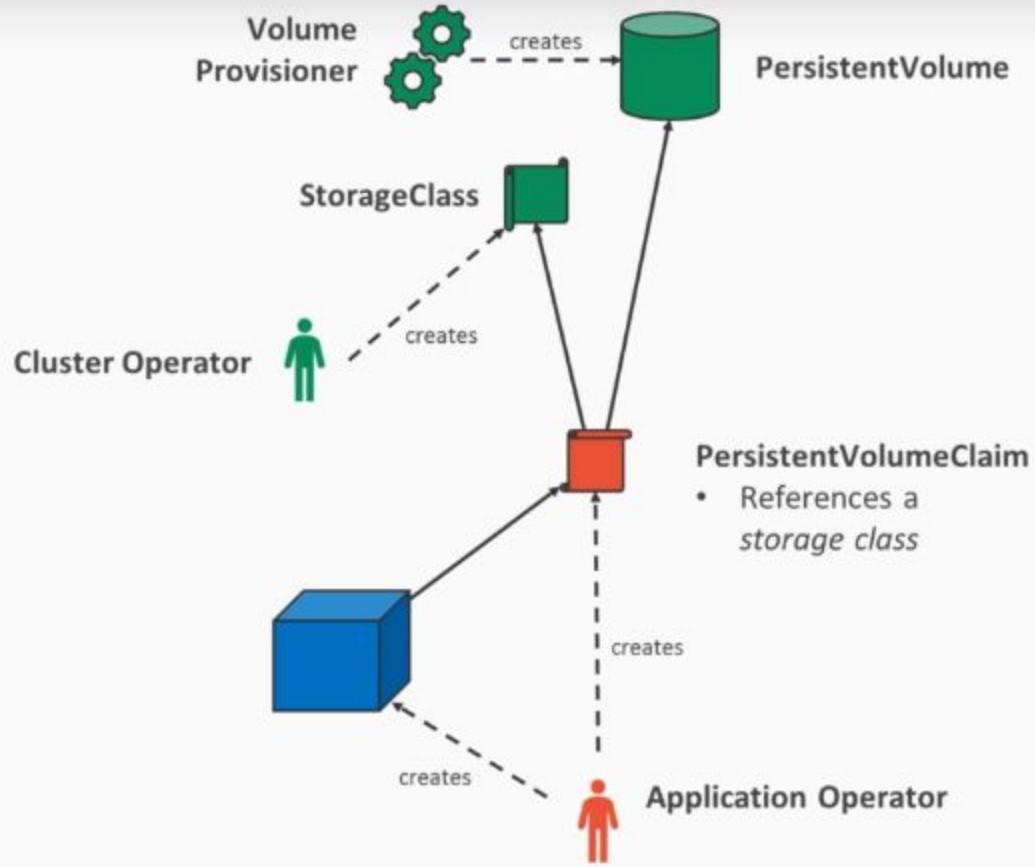
Storage Classes



Storage Classes



Storage Classes



StorageClass Example

```
apiVersion: storage.k8s.io/v1
```

```
kind: StorageClass
```

```
metadata:
```

```
    name: ssd-fast
```

```
provisioner: kubernetes.io/aws-ebs
```

```
parameters:
```

```
    type: io1
```

```
    zones: eu-central-1
```

```
    iopsPerGB: "20"
```

Provisioner name Examples:

- kubernetes.io/aws-ebs
- kubernetes.io/gce-pd
- kubernetes.io/azure-disk
- kubernetes.io/azure-file
- kubernetes.io/glusterfs

StorageClass Example

```
apiVersion: storage.k8s.io/v1
kind: StorageClass
metadata:
  name: ssd-fast
provisioner: kubernetes.io/aws-ebs
parameters:
  type: io1
  zones: eu-central-1
  iopsPerGB: "20"
Provisioner-dependent parameters
```

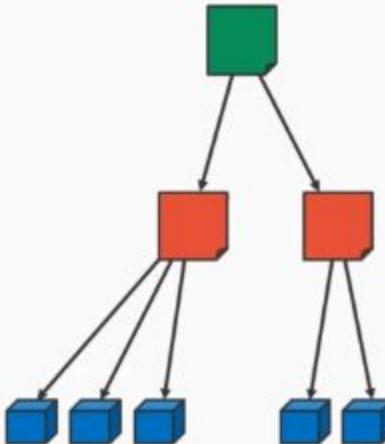
PersistentVolumeClaim Example with StorageClass

```
apiVersion: v1
kind: PersistentVolumeClaim
metadata:
  name: my-volume-claim
spec:
  + storageClass: ssd-fast
  resources:
    requests:
      storage: 10Gi
  accessModes:
    - ReadWriteOnce
```

```
~/Kubernetes/3.3-Automatic-Volume-Provisioning/Code> kubectl get pvc
NAME           STATUS    VOLUME                                     CAPACITY   ACCESS MODES  STORAGECLASS   AGE
my-volume-claim Bound    pvc-9f27911f-68f5-11e8-a87a-080027c553e8  5Gi       RWO          standard      17s
~/Kubernetes/3.3-Automatic-Volume-Provisioning/Code> kubectl get pv
NAME                                     CAPACITY   ACCESS MODES  RECLAIM POLICY  STATUS   CLAIM           STORAGECLASS   REASON   AGE
pvc-9f27911f-68f5-11e8-a87a-080027c553e8  5Gi       RWO          Delete        Bound    default/my-volume-claim standard   28s
~/Kubernetes/3.3-Automatic-Volume-Provisioning/Code>
```

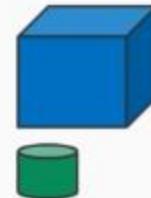
Using Stateful Sets

What Happened Previously?



Deployments and
RollingUpdates

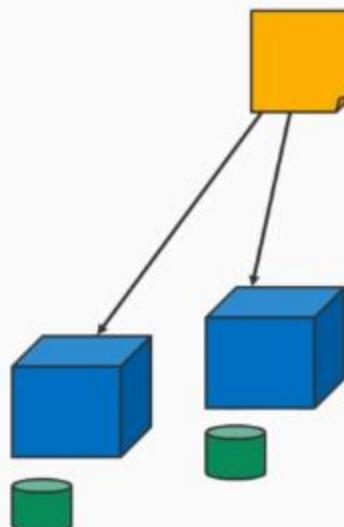
StatefulSets



Single Pod (stateful)

- Disposable
- Not resilient by itself
- Not scalable by itself

StatefulSets



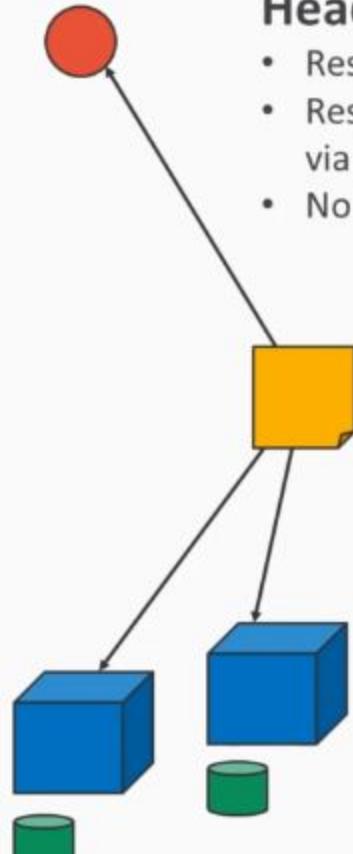
StatefulSet

- Manages Pods and volumes
- Self-healing and scalable
- Ideal for clustered apps

Single Pod (stateful)

- Disposable
- Not resilient by itself
- Not scalable by itself

StatefulSets



Headless Service

- Resolve StatefulSet via DNS
- Resolve individual members via DNS
- No IP loadbalancing, DNS only

StatefulSet

- Manages Pods and volumes
- Self-healing and scalable
- Ideal for clustered apps

Single Pod (stateful)

- Disposable
- Not resilient by itself
- Not scalable by itself

StatefulSet Example

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-database
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-database
  service: my-db-service
  template:
    metadata:
      labels:
        app: my-database
    spec:
      containers:
        - name: db
          image: mongo:3.7.9
          ports:
            # Omitted for brevity; see code samples
          volumeMounts:
            - name: data
              mountPath: /data
  volumeClaimTemplates:
    - metadata:
        name: data
      spec:
        accessModes: [ReadWriteOnce]
        storageClassName: ssd-fast
        resources:
          requests:
            storage: 10Gi
```

Replica count

StatefulSet Example

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-database
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-database
  service: my-db-service Assigned service name
  template:
    metadata:
      labels:
        app: my-database
    spec:
      containers:
        - name: db
          image: mongo:3.7.9
          ports:
            # Omitted for brevity; see code samples
          volumeMounts:
            - name: data
              mountPath: /data
      volumeClaimTemplates:
        - metadata:
            name: data
          spec:
            accessModes: [ReadWriteOnce]
            storageClassName: ssd-fast
            resources:
              requests:
                storage: 10Gi
```

StatefulSet Example

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-database
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-database
  service: my-db-service
  template:
    metadata:
      labels:
        app: my-database
    spec:
      containers:
        - name: db
          image: mongo:3.7.9
          ports:
            # Omitted for brevity; see code samples
      volumeMounts:
        - name: data
          mountPath: /data
  volumeClaimTemplates:
    - metadata:
        name: data
      spec:
        accessModes: [ReadWriteOnce]
        storageClassName: ssd-fast
        resources:
          requests:
            storage: 10Gi
```

Template for volume claims
(ideally with automatic provisioning)

StatefulSet Example

```
apiVersion: apps/v1
kind: StatefulSet
metadata:
  name: my-database
spec:
  replicas: 3
  selector:
    matchLabels:
      app: my-database
  service: my-db-service
  template:
    metadata:
      labels:
        app: my-database
    spec:
      containers:
        - name: db
          image: mongo:3.7.9
          ports:
            # Omitted for brevity; see code samples
      volumeMounts:
        - name: data
          mountPath: /data
volumeClaimTemplates:
  - metadata:
      name: data
    spec:
      accessModes: [ReadWriteOnce]
      storageClassName: ssd-fast
      resources:
        requests:
          storage: 10Gi
```

Using the volume claim in the Pod

Headless Services

```
apiVersion: v1
kind: Service
metadata:
  name: my-db-service
spec:
  ports:
    - port: 27017
      name: mongodb
  clusterIP: None
  selector:
    app: my-database
```

No IP load balancing;
DNS only

```
~/Kubernetes/3.4 Using StatefulSets/Code ➤ kubectl apply -f service.yaml
service "my-db-service" created
~/Kubernetes/3.4 Using StatefulSets/Code ➤ kubectl apply -f statefulset.yaml
statefulset.apps "my-database" created
~/Kubernetes/3.4 Using StatefulSets/Code ➤ kubectl get pods
NAME      READY   STATUS    RESTARTS   AGE
my-database-0  1/1     Running   0          7s
my-database-1  1/1     Running   0          6s
my-database-2  1/1     Running   0          5s
~/Kubernetes/3.4 Using StatefulSets/Code ➤ kubectl get pvc
NAME        STATUS    VOLUME                                     CAPACITY   ACCESS MODES  STORAGECLASS   AGE
data-my-database-0  Bound    pvc-30a0cbf7-6ca4-11e8-a8f0-080027c553e0  10Gi      RWO          standard       7h
data-my-database-1  Bound    pvc-41e97db0-6ca4-11e8-a8f0-080027c553e0  10Gi      RWO          standard       7h
data-my-database-2  Bound    pvc-4285cad1-6ca4-11e8-a8f0-080027c553e0  10Gi      RWO          standard       7h
my-volume-claim   Bound    pvc-9f27911f-68f5-11e8-a87a-080027c553e0  5Gi       RWO          standard       4d
~/Kubernetes/3.4 Using StatefulSets/Code ➤ kubectl run --rm -it --image alpine test
If you don't see a command prompt, try pressing enter.
!/ # nslookup my-db-service
nslookup: can't resolve '(null)': Name does not resolve

Name:      my-db-service
Address 1: 172.17.0.9 my-database-2.my-db-service.default.svc.cluster.local
Address 2: 172.17.0.7 my-database-0.my-db-service.default.svc.cluster.local
Address 3: 172.17.0.8 my-database-1.my-db-service.default.svc.cluster.local
/ #
```

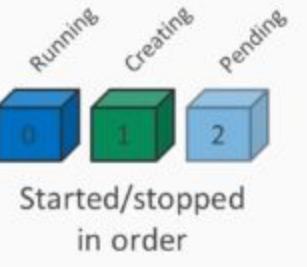
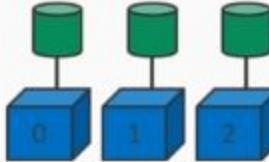
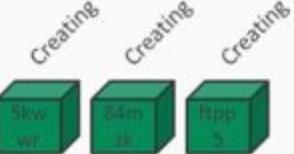
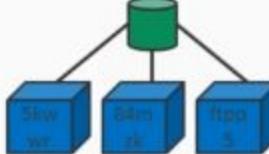
StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	Stable identity, DNS-addressable	Started/stopped in order	Volume provisioning for new members
ReplicaSet/ Deployment	No stable identity	Started/stopped all at once	Manually configured shared volumes

StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	Stable identity, DNS-addressable	Started/stopped in order	Volume provisioning for new members
ReplicaSet/ Deployment	No stable identity	Creating, Creating, Creating Started/stopped all at once	Manually configured shared volumes

StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	 <p>Stable identity, DNS-addressable</p>	 <p>Running Creating Pending</p> <p>Started/stopped in order</p>	 <p>Volume provisioning for new members</p>
ReplicaSet/ Deployment	 <p>No stable identity</p>	 <p>Creating Creating Creating</p> <p>Started/stopped all at once</p>	 <p>Manually configured shared volumes</p>

StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	Stable identity, DNS-addressable	Started/stopped in order	Volume provisioning for new members
ReplicaSet/ Deployment	No stable identity	Started/stopped all at once	Manually configured shared volumes

StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	Stable identity, DNS-addressable	Started/stopped in order	Volume provisioning for new members
ReplicaSet/ Deployment	No stable identity	Started/stopped all at once	Manually configured shared volumes

StatefulSets Versus Deployments

	Identity	Startup and shutdown order	Persistence
StatefulSet	Stable identity, DNS-addressable	Running, Creating, Pending Started/stopped in order	Volume provisioning for new members
ReplicaSet/ Deployment	No stable identity	Creating, Creating, Creating Started/stopped all at once	Manually configured shared volumes

StatefulSets Versus Deployments

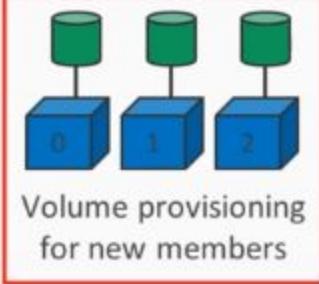
StatefulSet



Stable identity,
DNS-addressable



Started/stopped
in order

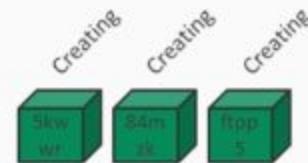


Volume provisioning
for new members

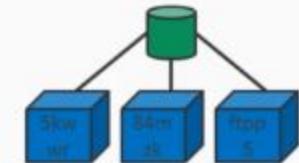
ReplicaSet/
Deployment



No stable identity



Started/stopped
all at once



Manually configured
shared volumes

Identity

Startup and
shutdown order

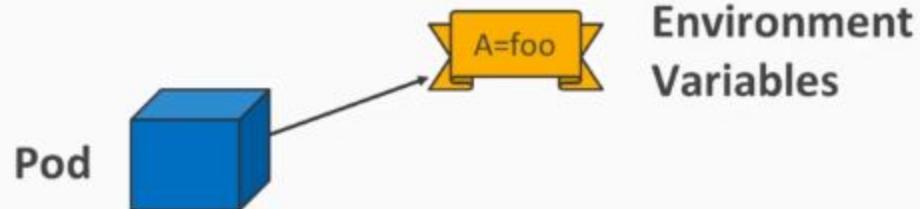
Persistence

Managing Configuration Data

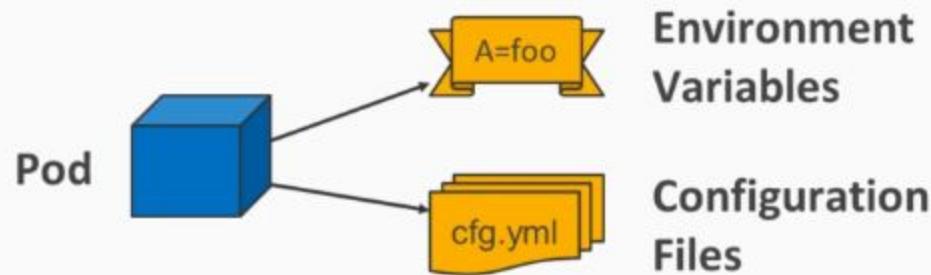
Pod Configuration



Pod Configuration



Pod Configuration



Pod Configuration

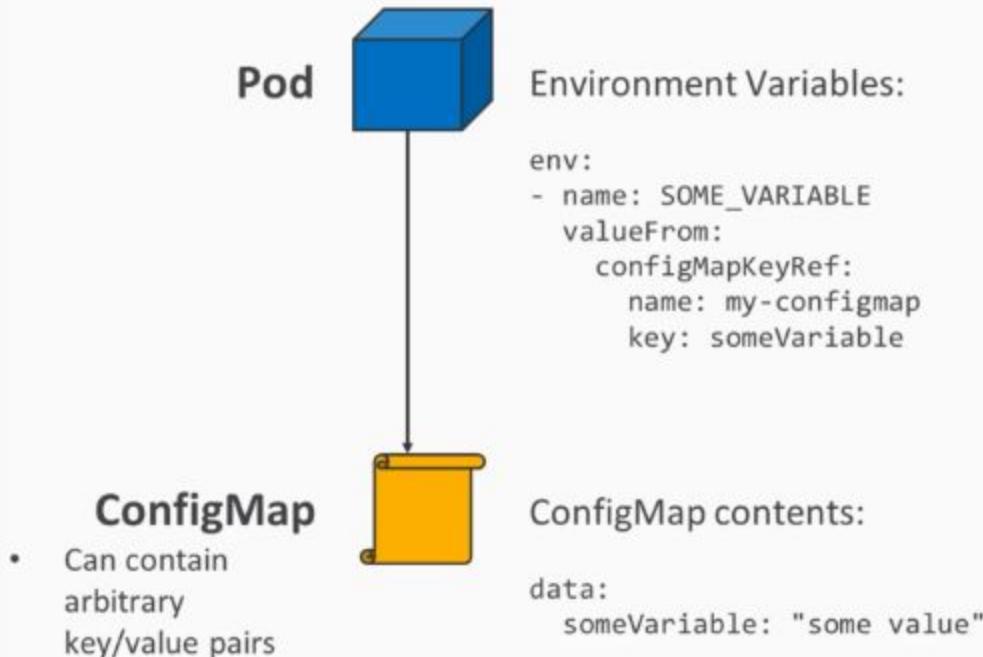
Pod



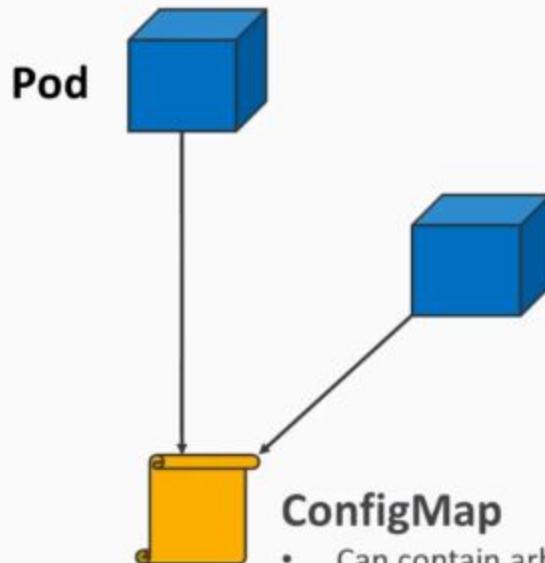
Environment Variables:

```
env:  
- name: SOME_VARIABLE  
  value: "some value"
```

Pod Configuration



Pod Configuration



ConfigMap

- Can contain arbitrary key/value pairs
- Can be used by multiple Pods at once
- Can be updated independently of the Pod

ConfigMap Example

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-configmap
data:
  someVariable: "some value"
  someBoolean: "true"
```

ConfigMap Example

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-configmap
data:
```

```
someVariable: "some value"
someBoolean: "true"
```

Key/value map

Caution:
Values must
be strings!

ConfigMap Usage Example

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
      env:
        - name: SOME_VARIABLE
          value: some value
        + valueFrom:
          + configMapKeyRef:
            + name: my-configmap
              key: someVariable
```

```
...a/4.1 Using ConfigMaps [master] ➔ kubectl apply -f configmap.yml  
configmap "my-configmap" created  
...a/4.1 Using ConfigMaps [master] ➔ kubectl get configmaps  
NAME      DATA   AGE  
my-configmap  2     3s  
...a/4.1 Using ConfigMaps [master]
```

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
  - name: app
    image: alpine
    command: ["sleep", "3600"]
    env:
    - name: SOME_VARIABLE
      valueFrom:
        configMapKeyRef:
          name: my-configmap
          key: someVariable
```

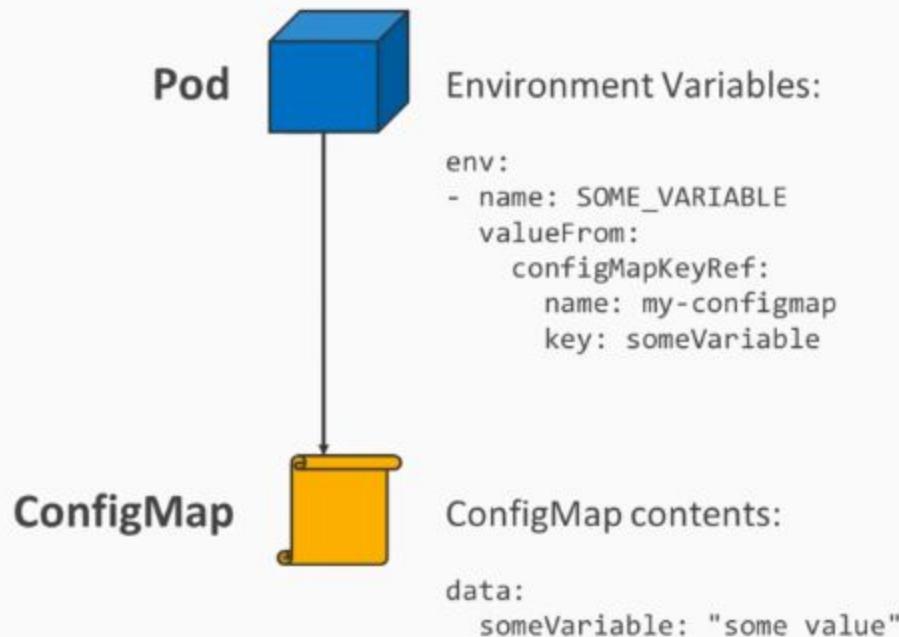
"pod.yaml" 15L, 267C

```
...a/4.1 Using ConfigMaps [master] kubectl apply -f configmap.yml  
configmap "my-configmap" created  
...a/4.1 Using ConfigMaps [master] kubectl get configmaps  
NAME      DATA   AGE  
my-configmap  2      3s  
...a/4.1 Using ConfigMaps [master] vim pod.yml  
...a/4.1 Using ConfigMaps [master]
```

```
...a/4.1 Using ConfigMaps $ master ➔ kubectl apply -f configmap.yml
configmap "my-configmap" created
...a/4.1 Using ConfigMaps $ master ➔ kubectl get configmaps
NAME      DATA   AGE
my-configmap   2      3s
...a/4.1 Using ConfigMaps $ master ➔ vim pod.yml
...a/4.1 Using ConfigMaps $ master ➔ kubectl apply -f pod.yml
pod "my-pod" created
...a/4.1 Using ConfigMaps $ master ➔ kubectl exec -it my-pod sh
/ # env
KUBERNETES_PORT=tcp://10.96.0.1:443
KUBERNETES_SERVICE_PORT=443
HOSTNAME=my-pod
SHLVL=1
HOME=/root
SOME_VARIABLE=some value
TERM=xterm
KUBERNETES_PORT_443_TCP_ADDR=10.96.0.1
PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
KUBERNETES_PORT_443_TCP_PORT=443
KUBERNETES_PORT_443_TCP_PROTO=tcp
KUBERNETES_PORT_443_TCP=tcp://10.96.0.1:443
KUBERNETES_SERVICE_PORT_HTTPS=443
PWD/
KUBERNETES_SERVICE_HOST=10.96.0.1
/ # █
```

Inject Configuration Files Pods

Previously



ConfigMap Example

```
apiVersion: v1
kind: ConfigMap
metadata:
  name: my-configmap
data:
  config.ini: |  YAML multiline string
    some_variable = foo
    some_other_variable = baz
  database.ini: |
    uri = mongodb://my-
database:27017/...
```

ConfigMap Usage Example

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
    +  volumeMounts:
    +    - name: config
    +      mountPath: /etc/config
    +  volumes:
    +    - name: config
    +      configMap:
    +        name: my-configmap
```

ConfigMap Usage Example

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
      + volumeMounts:
        - name: config
          mountPath: /etc/config
      + volumes:
        - name: config
      + configMap:
        + name: my-configmap
```

This mount is automatically updated when the ConfigMap changes!

```
...ng Configuration Files | master | cd config  
...figuration Files/config | master | ls -l  
total 16  
-rw-r--r-- 1 mhelmich staff 46 11 Jun 19:28 config.ini  
-rw-r--r-- 1 mhelmich staff 39 11 Jun 19:28 database.ini  
...figuration Files/config | master | kubectl create configmap my-configmap --from-file=
```

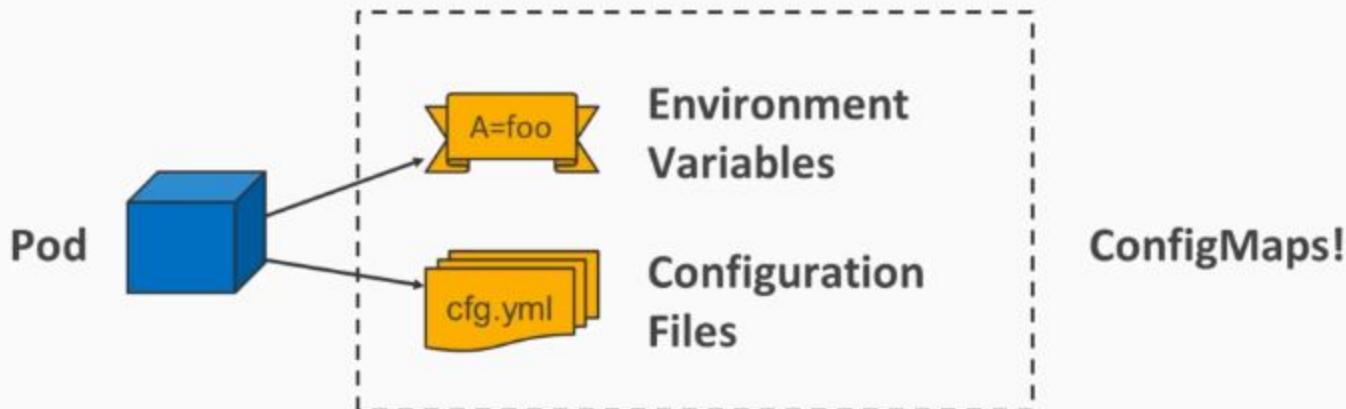
```
■ Please edit the object below. Lines beginning with a '#' will be ignored,  
# and an empty file will abort the edit. If an error occurs while saving this file will be  
# reopened with the relevant failures.  
#  
apiVersion: v1  
data:  
config.ini: |  
    some_variable = foo  
    some_other_variable = baz  
database.ini: |  
    uri = mongodb://my-database:27017/test  
kind: ConfigMap  
metadata:  
creationTimestamp: 2018-06-12T21:10:52Z  
name: my-configmap  
namespace: default  
resourceVersion: "58518"  
selflink: /api/v1/namespaces/default/configmaps/my-configmap  
uid: 1698f66b-6e85-11e8-b829-080027c553e0
```

```
"/var/folders/lz/z3f_tz6953jchlbkswp0zy00000gn/T/kubectl-edit-llczx.yaml" 19L, 604C
```

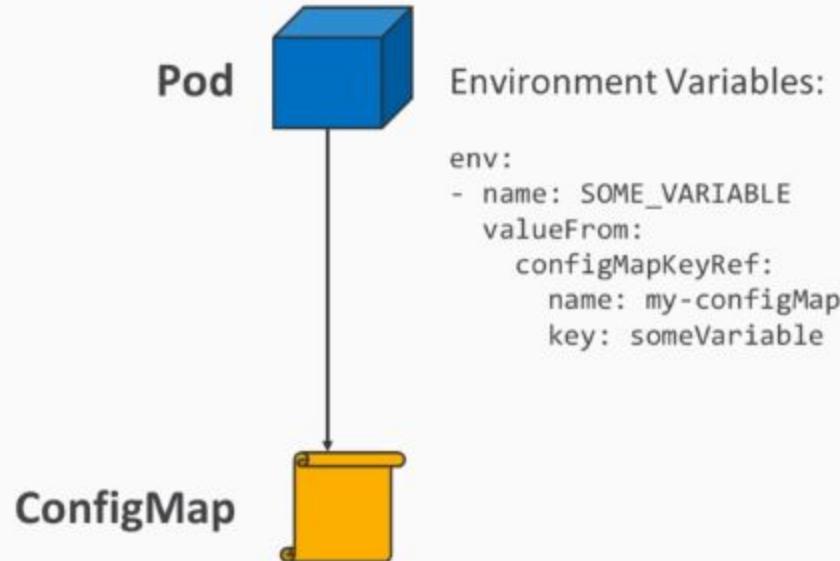
```
...figuration Files/config      [master] cd ..
...ng Configuration Files      [master] kubectl apply -F pod.yml
pod "my-pod" created
...ng Configuration Files      [master] kubectl exec -it my-pod sh
/ # ls -l /etc/config
total 0
lrwxrwxrwx  1 root    root          17 Jun 12 21:12 config.ini -> ..data/config.ini
lrwxrwxrwx  1 root    root          19 Jun 12 21:12 database.ini -> ..data/database.ini
/ # cat /etc/config/config.ini
some_variable = foo
some_other_variable = baz
/ #
```

Keeping Secrets

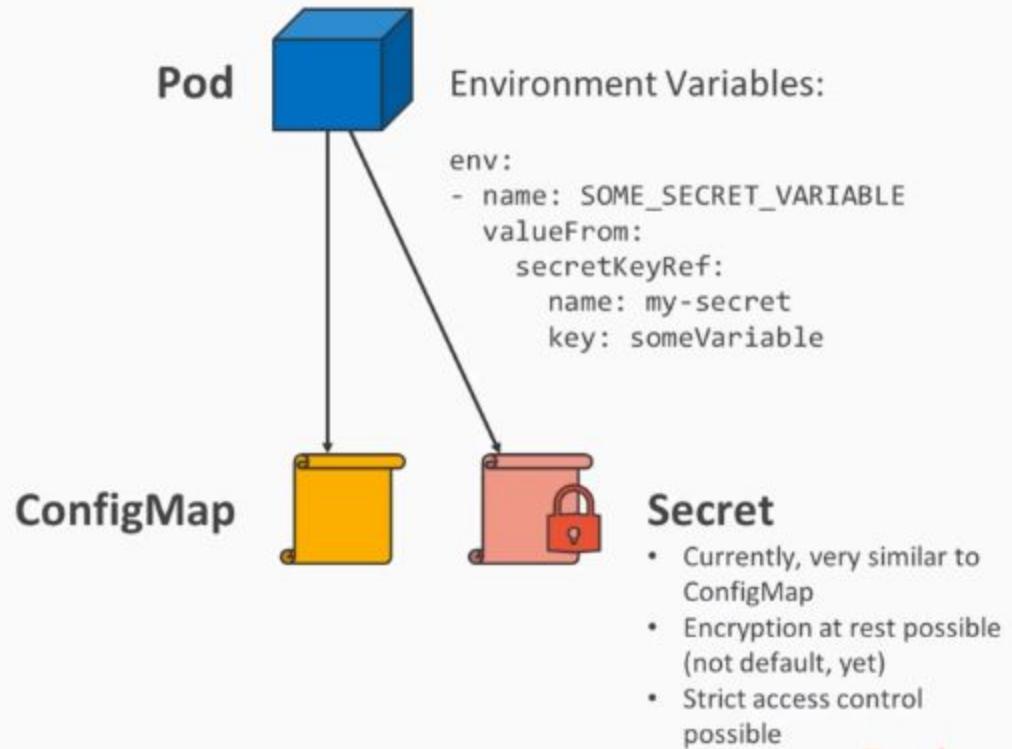
Previously



Pod Secrets



Pod Secrets



Secret Example

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
data:
  test: VGVzdA==
```

Secret Example

```
apiVersion: v1
kind: Secret
metadata:
  name: my-secret
data:
  test: VGVzdA==
```

Arbitrary key/value pairs
Values are base64 encoded

base64 is not encryption!
base64 is not encryption!

Secret Usage in Environment Variables

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
    + env:
      + - name: SOME_SECRET_VARIABLE
        valueFrom:
          secretKeyRef:
            + name: my-secret
            + key: someVariable
```

Secret Usage in Volumes

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
    +  volumeMounts:
    +    - name: secret
    +      mountPath: /etc/secret
    +  volumes:
    +    - name: secret
    +      secret:
    +        name: my-secret
```

...ta/4.3 Keeping Secrets master

I

```
...ta/4.3 Keeping Secrets ✘ master ➔ kubectl create secret generic my-secret --from-file=.
secret "my-secret" created
...ta/4.3 Keeping Secrets ✘ master ➔
```

```
...ta/4.3 Keeping Secrets ✘ master ➤ kubectl create secret generic my-secret --from-file=.
secret "my-secret" created
...ta/4.3 Keeping Secrets ✘ master ➤ ls -l
total 8
-rw-r--r-- 1 mhelmich staff 41 14 Jun 21:43 credentials.yaml
...ta/4.3 Keeping Secrets ✘ master ➤ kubectl edit secret my-secret
Edit cancelled, no changes made.
...ta/4.3 Keeping Secrets ✘ master ➤
```

Composing Environment Variables

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
      env:
        - name: MONGODB_PASSWORD
          valueFrom:
            secretKeyRef:
              name: mongodb-credentials
              key: password
        - name: MONGODB_URL
          value: "mongodb://user:$(MONGODB_PASSWORD)@mongo/database"
```

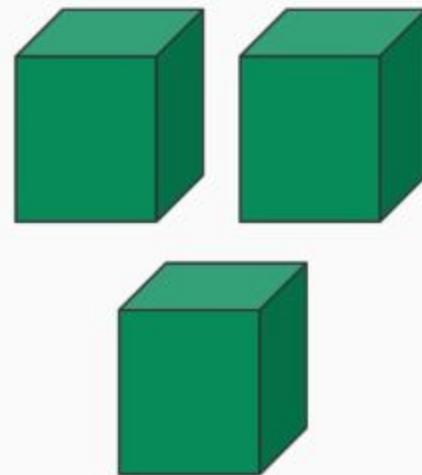
Composing Environment Variables

```
apiVersion: v1
kind: Pod
metadata:
  name: my-pod
spec:
  containers:
    - name: my-service
      image: nginx:1.3.12
      env:
        - name: MONGODB_PASSWORD
          valueFrom:
            secretKeyRef:
              name: mongodb-credentials
              key: password
        - name: MONGODB_URL
          value: "mongodb://user:$({{MONGODB_PASSWORD}}@mongo/database")"
```

Continuous Delivery with Kubernetes

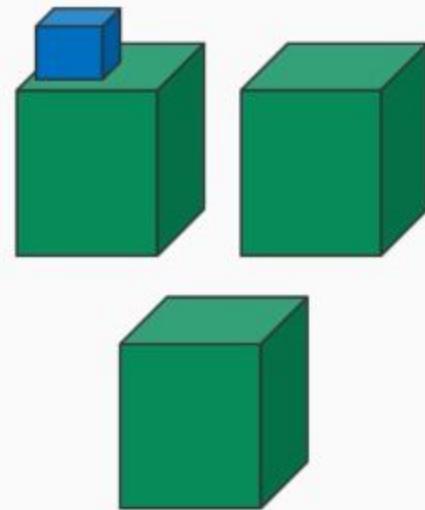
Characteristics of a CI/CD Pipeline

Continuous What?



Your Kubernetes
Cluster

Continuous What?

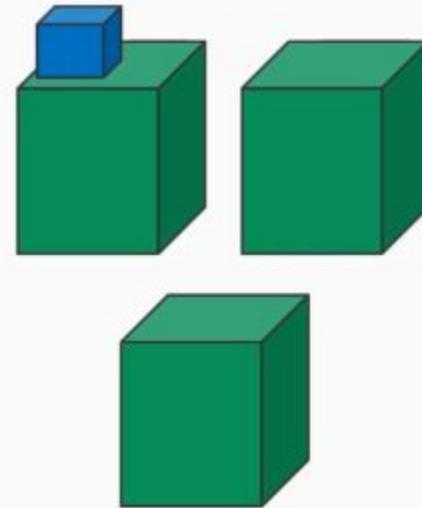


Your Kubernetes
Cluster

Continuous What?



Your Source Code
Repository



Your Kubernetes
Cluster

A Typical Build Pipeline

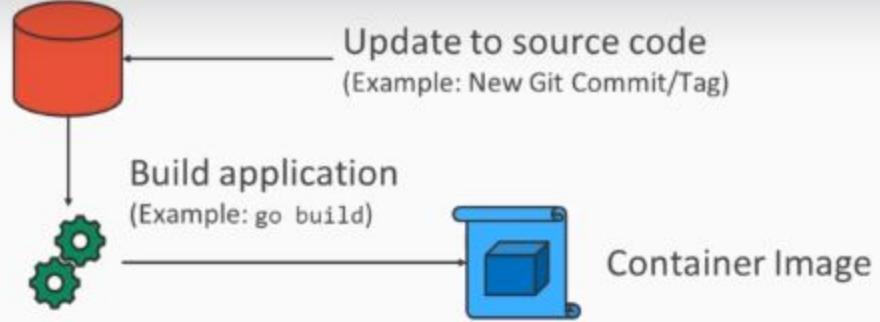


Update to source code
(Example: New Git Commit/Tag)

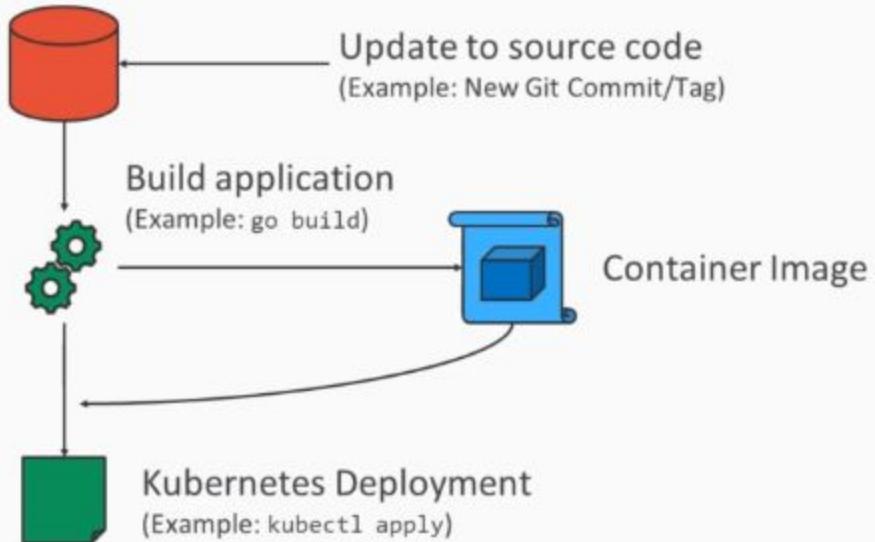
A Typical Build Pipeline



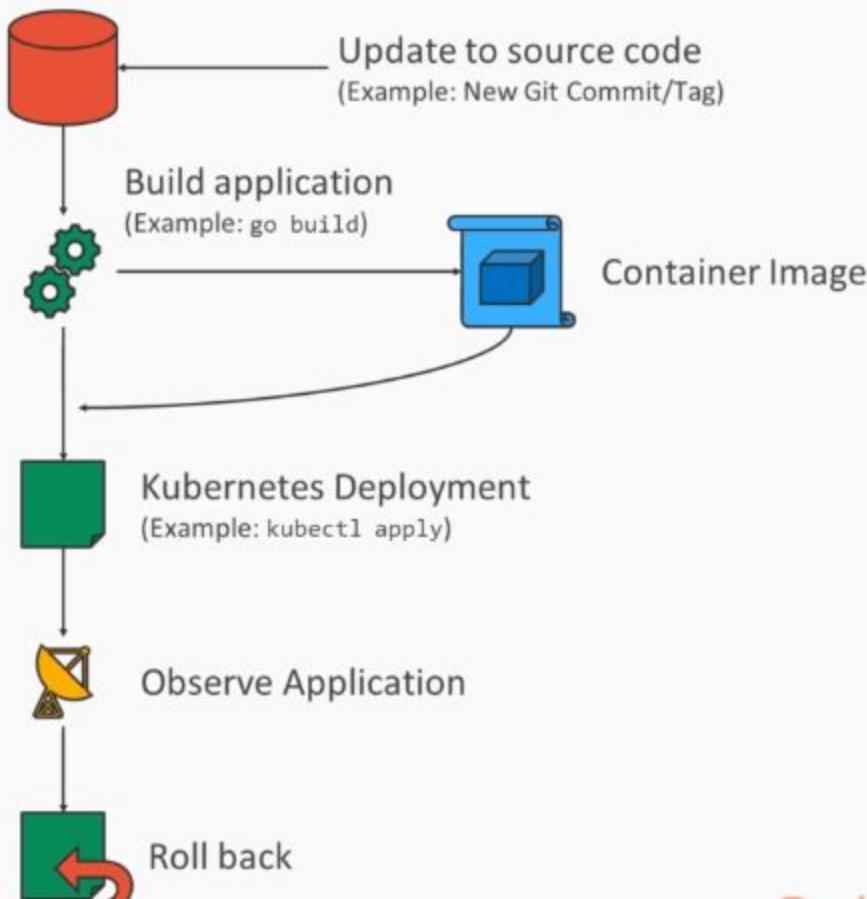
A Typical Build Pipeline



A Typical Build Pipeline



A Typical Build Pipeline



How to Release a Specific Version?



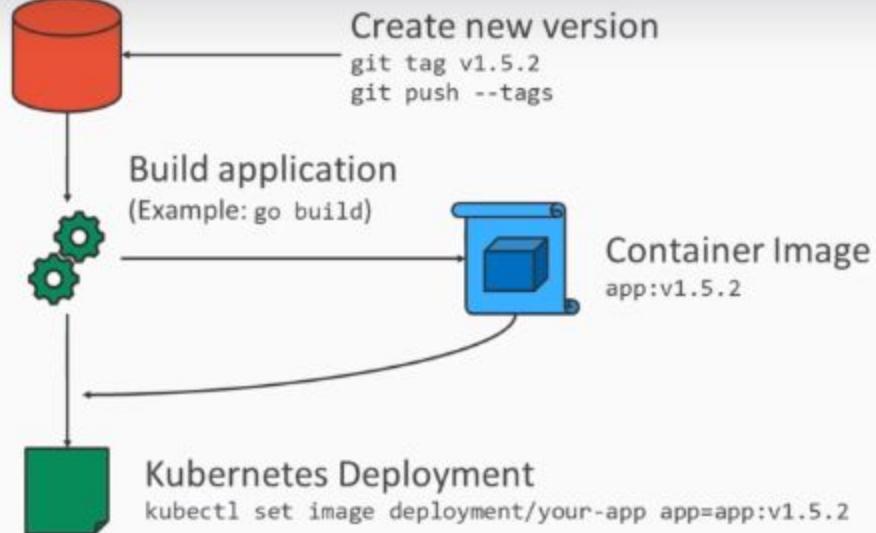
Create new version

```
git tag v1.5.2  
git push --tags
```

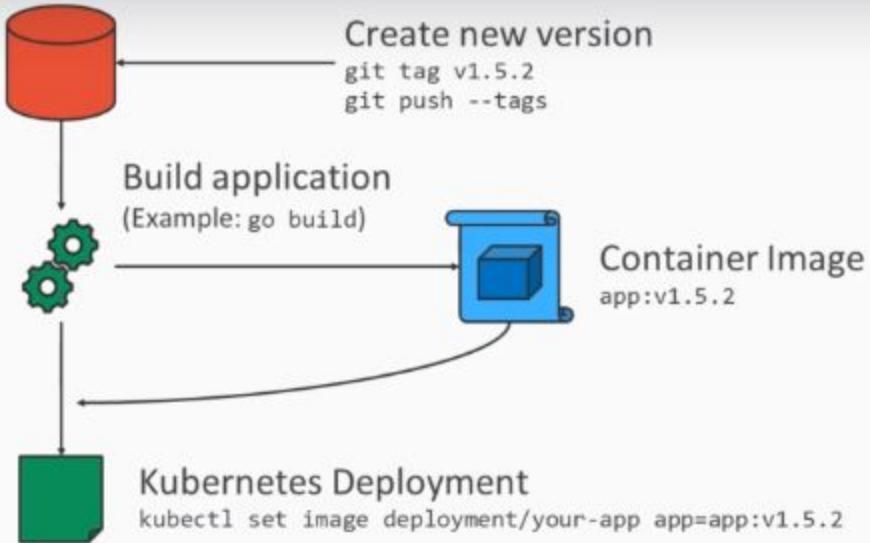
How to Release a Specific Version?



How to Release a Specific Version?



How to Release a Specific Version?



Caution!

- Using `set image` requires that Deployment already exists
- What about Updates to other parts of the Pod spec (environment variables, ports, and so on)?

Later in This Section



GitLab



Configuring CD with GitLab



Why
GitLab ?

Prerequisites

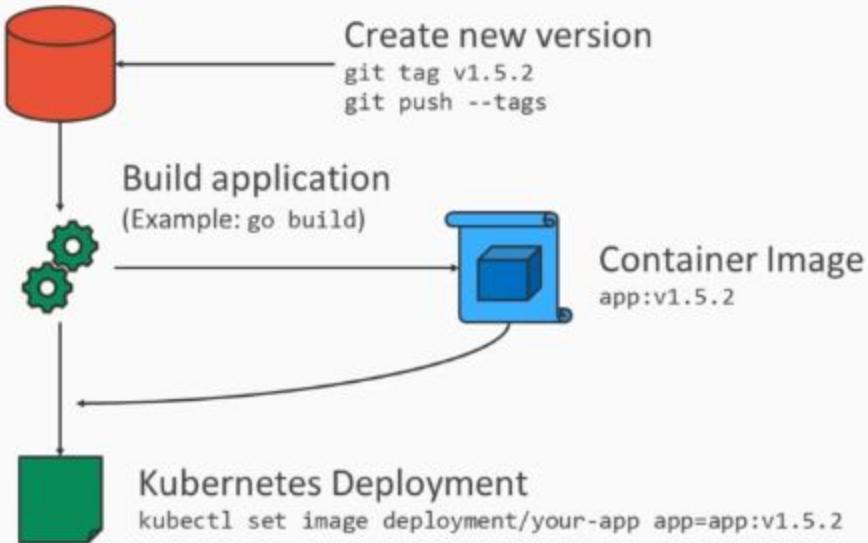
- Running GitLab instance with CI runner
(alternatively: hosted GitLab at
gitlab.com)
- Running Kubernetes cluster, reachable
from CI runner
- Alternatively: Have a look at the
example files for this video for a quick
and easy (but not production-ready)
GitLab setup on Minikube



Set Up GitLab on Kubernetes

[https://docs.gitlab.com/ce/install/
kubernetes/](https://docs.gitlab.com/ce/install/kubernetes/)

The Goal



```
my-project/
  .gitignore
  .gitlab-ci.yml
  Dockerfile
  main.go
```

```
my-project/
  .gitignore
  .gitlab-ci.yml
  Dockerfile
  main.go
```

my-project/
 .gitignore
 .gitlab-ci.yml
 Dockerfile
 main.go

```
stages: [build, package, deploy]

compile:
  stage: build
  # ... omitted for brevity

dockerbuild:
  stage: package
  image: docker:18.05.0-ce
  script:
    - docker build -t app:${CI_COMMIT_TAG} .
    - docker push app:${CI_COMMIT_TAG}
  only:
    - tags

deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - >
      kubectl set image deployment/your-app
      app=app:${CI_COMMIT_TAG}
    - kubectl get pods
  only:
    - tags
```

my-project/
 .gitignore
 .gitlab-ci.yml
 Dockerfile
 main.go

```
stages: [build, package, deploy]

compile:
  stage: build
  # ... omitted for brevity

dockerbuild:
  stage: package
  image: docker:18.05.0-ce
  script:
    - docker build -t app:${CI_COMMIT_TAG} .
    - docker push app:${CI_COMMIT_TAG}
  only:
    - tags

deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - >
      kubectl set image deployment/your-app
      app=app:${CI_COMMIT_TAG}
    - kubectl get pods
  only:
    - tags
```

my-project/
 .gitignore
 .gitlab-ci.yml
 Dockerfile
 main.go

```
stages: [build, package, deploy]

compile:
  stage: build
  # ... omitted for brevity

dockerbuild:
  stage: package
  image: docker:18.05.0-ce
  script:
    - docker build -t app:${CI_COMMIT_TAG} .
    - docker push app:${CI_COMMIT_TAG}
  only:
    - tags

deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - >
      kubectl set image deployment/your-app
      app=app:${CI_COMMIT_TAG}
    - kubectl get pods
  only:
    - tags
```

my-project/
 .gitignore
 .gitlab-ci.yml
 Dockerfile
 main.go

```
stages: [build, package, deploy]

compile:
  stage: build
  # ... omitted for brevity See documentation for more information
  # https://docs.gitlab.com/ce/ci/docker/using_docker_build.html

  dockerbuild:
    stage: package
    image: docker:18.05.0-ce
    script:
      - docker build -t app:${CI_COMMIT_TAG} .
      - docker push app:${CI_COMMIT_TAG}
    only:
      - tags

  deploy:
    stage: deploy
    image: quay.io/martinhelmich/k8s-deployer:v1.0.0
    environment: development
    script:
      - >
        kubectl set image deployment/your-app
        app=app:${CI_COMMIT_TAG}
      - kubectl get pods
    only:
      - tags
```

my-project/
.gitignore
.gitlab-ci.yml
Dockerfile
main.go

```
stages: [build, package, deploy]

compile:
  stage: build
  # ... omitted for brevity

dockerbuild:
  stage: package
  image: docker:18.05.0-ce
  script:
    - docker build -t app:${CI_COMMIT_TAG} .
    - docker push app:${CI_COMMIT_TAG}
  only:
    - tags

deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - >
      kubectl set image deployment/your-app
      app=app:${CI_COMMIT_TAG}
    - kubectl get pods
  only:
    - tags
```

192.168.99.100:30080/mhelmich/test

GitLab

Projects Groups Activity Milestones Snippets

This project Search

test

Martin Helmich > test > Details

T test

Project Details Activity Cycle Analytics

Star 0 Fork 0 SSH git@gitlab-99bb7bd74-bj9pn:mhe

Global

Repository Issues Merge Requests CI / CD Operations Wiki Snippets Settings

Files (2.5 MB) Commits (4) Branch (1) Tags (0) CI/CD configuration

Add Changelog Add License Add Contribution guide Add Kubernetes cluster

master / test / + History Find file Web IDE

Update CI config Martin Helmich authored 9 minutes ago 88bcc951

Name	Last commit	Last update
.gitlab-ci.yml	Update CI config	9 minutes ago

GitLab Projects Groups Activity Milestones Snippets 🔍

+ This project Search 🔍 🌐

T test

Project Details Activity Cycle Analytics Repository Issues Merge Requests CI / CD Operations Wiki Snippets

Files (2.5 MB) Commits (4) Branch (1) Tags (0) CI/CD configuration

Add Changelog Add License Add Contribution guide Add Kubernetes cluster

master test / + History Find file Web IDE ⚙️

Update CI config Martin Helmich authored 9 minutes ago 88bcc951 📁

Name	Last commit	Last update
.gitlab-ci.yml	Update CI config	9 minutes ago
Dockerfile	Add dockerfile and fix pipeline	1 day ago
deployment.yaml	Update CI config	9 minutes ago
example-app	Initial commit	1 day ago
main.go	Initial commit	1 day ago

GitLab Projects Groups Activity Milestones Snippets 🔍

Martin Helmich > test > Repository

master test / Dockerfile Find file Blame History Permalink

Add dockerfile and fix pipeline
Martin Helmich authored 1 day ago

Dockerfile 78 Bytes Edit Web IDE Replace Delete

```
1 FROM scratch
2
3 COPY example-app /example-app
4
5 EXPOSE 8080
6
7 CMD ["/example-app"]
```

T test Project Repository Files Commits Branches Tags Contributors Graph Compare Charts Issues 0 Merge Requests 0 CI / CD

GitLab Projects Groups Activity Milestones Snippets

This project Search

T test

Project Repository

Files Commits Branches Tags Contributors Graph Compare Charts

Issues 0 Merge Requests 0 CI / CD

Collapse sidebar

.gitlab-ci.yml 768 Bytes

```
stages:
  - build
  - package
  - deploy

compile:
  stage: build
  image: golang:1.10.3
  script:
    - go build -o example-app
  artifacts:
    paths:
      - example-app
    expire_in: 1 week

dockerbuild:
  stage: package
  image: docker:18.05.0-ce
  script:
    - docker build -t app:${CI_COMMIT_TAG} .
    - docker tag app:${CI_COMMIT_TAG} app:latest

# Disabled as long as we're running on minikube
# Since it's all just one VM, any built image will be available immediately, anyway
#- docker push app:${CI_COMMIT_TAG}
#- docker push app:latest

only:
  - tags
```

```
...@1th: GitLab/example-app ➜ master ➜ █
```

GitLab

Projects Groups Activity Milestones Snippets

This project Search

Project Repository Issues Merge Requests

CI / CD Pipelines Jobs Schedules Charts Operations Wiki Snippets

test

Martin Helmich > test > Pipelines

All 5 Pending 0 Running 1 Finished 4 Branches Tags

Run Pipeline Clear Runner Caches CI Lint

Status	Pipeline	Commit	Stages	
running	#5 by 🎉	v1.0.1 -> 88bcc951 Update CI config	🟡 ⚡ ⚡	X
passed	#4 by 🎉	master -> 88bcc951 Update CI config	✓	00:00:11 11 minutes ago
passed	#3 by 🎉	master -> a04113d1 Add dockerfile and fix pi...	✓	00:00:09 1 day ago
passed	#2 by 🎉	master -> 7c58202f Fix CI	✓ ✓	00:00:42 1 day ago
failed	#1 by 🎉	master -> 95a387f5 Initial commit		1 day ago

 GitLab Projects Groups Activity Milestones Snippets 🔍

This project Search 🔎 🗂️ 🌐 📄

T test Martin Helmich > test > Pipelines > #5

 Pipeline #5 triggered 1 minute ago by  Martin Helmich

Update CI config

⌚ 3 jobs from v1.0.1 in 1 minute 2 seconds

→ 88bcc951 ⏪

Pipeline Jobs 3

Build	Package	Deploy			
 compile		 dockerbuild		 deploy	

🕒 CI / CD

Project Repository Issues Merge Requests CI / CD Pipelines Jobs Schedules Charts Operations Wiki Snippets

192.168.99.100:30060/mhelmich/test/-/jobs/7

GitLab

Projects Groups Activity Milestones Snippets

This project Search

T test

Project Repository Issues Merge Requests CI / CD Pipelines Jobs Schedules Charts Operations Wiki Snippets

Running
Waiting for pod default/runner-5eeb8561-project-1-concurrent-0czwx to be running, status is Pending
Waiting for pod default/runner-5eeb8561-project-1-concurrent-0czwx to be running, status is Pending
Running on runner-5eeb8561-project-1-concurrent-0czwx via gitlab-runner-96bd68669-b9xdx...
Cloning repository...
Cloning into '/mhelmich/test'...
Checking out 88bcc951 as v1.0.1...
Skipping Git submodules setup
Downloading artifacts for compile (6)...
Downloading artifacts from coordinator... ok id=6 responseStatus=200 OK token=6znvR6xp
\$ docker build -t app:\${CI_COMMIT_TAG} .
Sending build context to Docker daemon 9.064MB

Step 1/4 : FROM scratch
-->
Step 2/4 : COPY example-app /example-app
--> f8729564ee42
Step 3/4 : EXPOSE 8080
--> Running in 81f8874ebb92
--> a076676b89be
Removing intermediate container 81f8874ebb92
Step 4/4 : CMD /example-app
--> Running in 7a62aa8cb3a8
--> 8b6be8a7f81c
Removing intermediate container 7a62aa8cb3a8
Successfully built 8b6be8a7f81c
Successfully tagged app:v1.0.1
\$ docker tag app:\${CI_COMMIT_TAG} app:latest
Job succeeded

dockerbuild

Retry

Duration: 38 seconds
Timeout: 1h (from project)
Runner: gitlab-runner-96bd68669-xrpzf (#1)

Commit 88bcc951

Update CI config

Pipeline #5 from v1.0.1

deploy

deploy

GitLab Projects Groups Activity Milestones Snippets + This project Search

T test Martin Helmich > test > Jobs > #8

Project Repository Issues Merge Requests CI / CD Pipelines Jobs Schedules Charts Operations Wiki Snippets

passed Job #8 triggered 52 seconds ago by Martin Helmich

```
Running with gitlab-runner 11.0.0 (5396d320)
on gitlab-runner-96bd68669-xrpzf Seeb8561
WARNING: Namespace is empty, therefore assuming 'default'.
Using Kubernetes namespace: default
Using Kubernetes executor with image quay.io/martinhelmich/k8s-deployer:v1.0.0 ...
Waiting for pod default/runnner-5eeb8561-project-1-concurrent-05mf97 to be running, status is Pending
Waiting for pod default/runnner-5eeb8561-project-1-concurrent-05mf97 to be running, status is Pending
Running on runner-5eeb8561-project-1-concurrent-05mf97 via gitlab-runner-96bd68669-b9xdx...
Cloning repository...
Cloning into '/mhelmich/test'...
Checking out 88bcc951 as v1.0.1...
Skipping Git submodules setup
Downloading artifacts for compile (6)...
Downloading artifacts from coordinator... ok      id=6 responseStatus=200 OK token=6znvR6xp
$ kubectl set image deployment/my-deployment app=app:${CI_COMMIT_TAG}
deployment.apps "my-deployment" image updated
Job succeeded
```

deploy Retry

Duration: 13 seconds
Timeout: 1h (from project)
Runner: gitlab-runner-96bd68669-xrpzf (#1)

Commit 88bcc951

Update CI config

Pipeline #5 from v1.0.1

deploy

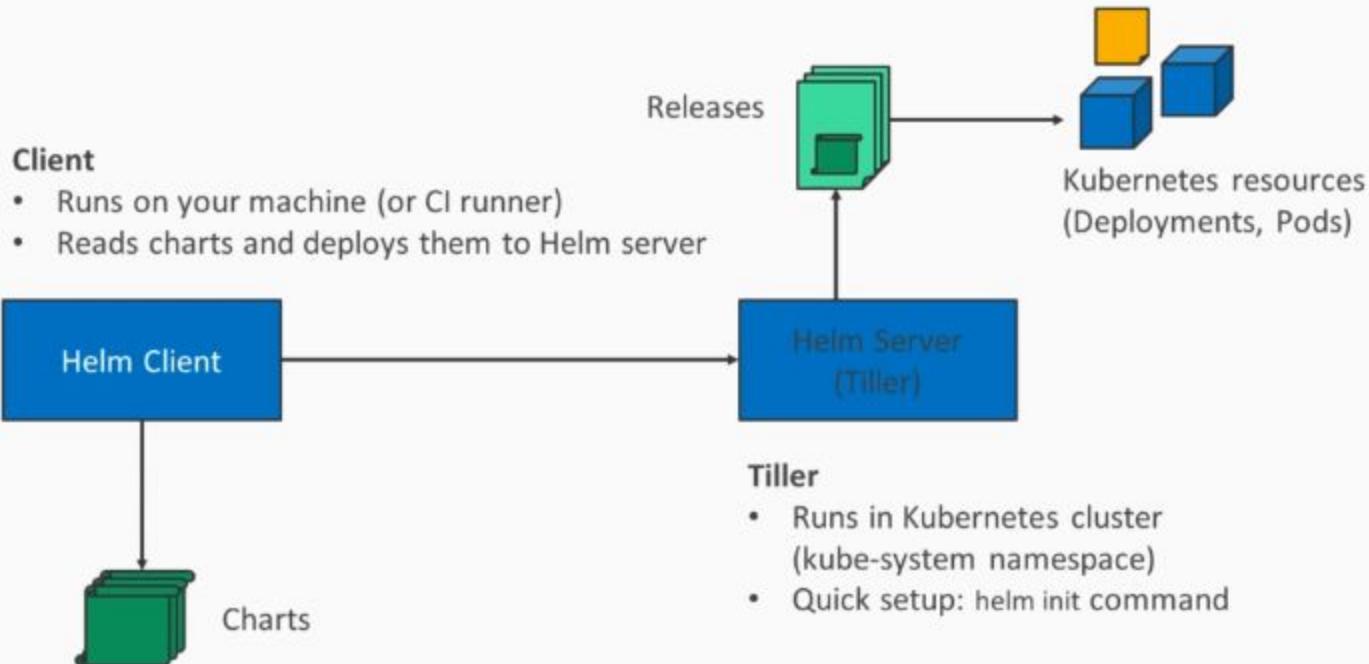
→ deploy

```
...@...:~/GitLab/example-app> git tag v1.0.1
...@...:~/GitLab/example-app> git push --tags
warning: redirecting to http://192.168.99.100:30080/mhelmich/test.git/
Total 0 (delta 0), reused 0 (delta 0)
To http://192.168.99.100:30080/mhelmich/test
 * [new tag]      v1.0.1 -> v1.0.1
...@...:~/GitLab/example-app> git status
```

```
...ith Gitlab/example-app ➜ master ➜ git tag v1.0.1
...ith Gitlab/example-app ➜ master ➜ git push --tags
warning: redirecting to http://192.168.99.100:30080/mhelmich/test.git/
Total 0 (delta 0), reused 0 (delta 0)
To http://192.168.99.100:30080/mhelmich/test
 * [new tag]      v1.0.1 -> v1.0.1
...ith Gitlab/example-app ➜ master ➜ kubectl get pods
NAME          READY   STATUS    RESTARTS   AGE
gitlab-99bb7bd74-bj9pn   1/1     Running   1          1d
gitlab-runner-96bd68669-b9wdx   1/1     Running   1          1d
my-deployment-6b66998c44-xrn26   0/1     CrashLoopBackOff   3          1m
my-deployment-d99bb7895-6hxnn   1/1     Running   0          16m
my-deployment-d99bb7895-9t49k   1/1     Running   0          16m
my-deployment-d99bb7895-siblx   1/1     Running   0          16m
...ith Gitlab/example-app ➜ master ➜
```

Setting up Helm

Helm Architecture



Helm Installation and Setup

1. Download and install Helm to your local machine
https://docs.helm.sh/using_helm/#installing-helm
2. Setup local Helm and Tiller server using
`$ helm init`

```
...ng up Helm/example-app [master] ➜ helm init  
$HELM_HOME has been configured at /Users/mhelmich/.helm.  
  
Tiller (the Helm server-side component) has been installed into your Kubernetes Cluster.  
Happy Helming!  
...ng up Helm/example-app [master] ➜
```

Helm Chart Structure

```
my-chart/
  Chart.yaml
  values.yaml
  templates/
    deployment.yaml
    service.yaml
```

Helm Chart Structure

```
my-chart/  
  Chart.yaml  
  values.yaml  
  templates/  
    deployment.yaml  
    service.yaml
```

```
apiVersion: v1  
description: A Helm chart for your application  
name: my-app  
version: 0.1.0
```

Helm Chart Structure

```
my-chart/
  Chart.yaml
  values.yaml
  templates/
    deployment.yaml
    service.yaml
```

```
replicaCount: 1
image:
  repository: app
  tag: latest
  pullPolicy: Always
service:
  name: http
  type: ClusterIP
  externalPort: 80
  internalPort: 80
```

Helm Chart Structure

my-chart/

 Chart.yaml

 values.yaml

 templates/

 deployment.yaml

 service.yaml

```
apiVersion: apps/v1
kind: Deployment
metadata:
  name: {{ template "chart.fullname" . }}
  labels:
    app: {{ template "chart.name" . }}
    chart: {{ .Chart.Name }}-{{ .Chart.Version | replace "+" "_" }}
    release: {{ .Release.Name }}
    heritage: {{ .Release.Service }}
spec:
  replicas: {{ .Values.replicaCount }}
  template:
    metadata:
      labels:
        app: {{ template "chart.name" . }}
        release: {{ .Release.Name }}
    spec:
      containers:
        - name: {{ .Chart.Name }}
          image: {{ .Values.image.repository }}:{{ .Values.image.tag }}
          imagePullPolicy: {{ .Values.image.pullPolicy }}
        ports:
          - protocol: TCP
            name: http
            containerPort: {{ .Values.service.internalPort }}
```

Helm Chart Structure

```
my-chart/  
  Chart.yaml  
  values.yaml  
  templates/  
    deployment.yaml  
    service.yaml
```

Create these files for a new project quickly using the helm create <name> command

```
...ng up Helm/example-app [master] ➔ ls -al
total 11968
drwxr-xr-x  8 mhelmich  staff    256 Jun 13:21 .
drwxr-xr-x  3 mhelmich  staff     96 Jun 13:21 ..
drwxr-xr-x 12 mhelmich  staff   384 Jun 14:56 .git
-rw-r--r--  1 mhelmich  staff   768 Jun 14:39 .gitlab-ci.yml
-rw-r--r--  1 mhelmich  staff    78 Jun 11:25 Dockerfile
drwxr-xr-x  7 mhelmich  staff   224 Jun 13:21 chart
-rwxr-xr-x  1 mhelmich  staff  6110461 Jun 10:19 example-app
-rw-r--r--  1 mhelmich  staff   198 Jun 10:18 main.go
...ng up Helm/example-app [master] ➔ helm install --set image.tag=v1.0.0 --name myrelease ./chart
NAME: myrelease
LAST DEPLOYED: Sun Jun 24 14:58:01 2018
NAMESPACE: default
STATUS: DEPLOYED

RESOURCES:
==> v1/Service
NAME          TYPE      CLUSTER-IP      EXTERNAL-IP      PORT(S)      AGE
myrelease-chart ClusterIP  10.102.11.62 <none>        80/TCP      1s

==> v1beta1/Deployment
NAME          DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
myrelease-chart 1         1         1           0          1s

==> v1/Pod(related)
NAME          READY  STATUS          RESTARTS  AGE
myrelease-chart-5755b87555-qjc25  0/1  ContainerCreating  0          1s

NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app=chart,release=myrelease" -o jsonpath=".items[0].metadata.name")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl port-forward $POD_NAME 8080:80
```

```

==> v1/Pod(related)
NAME READY STATUS RESTARTS AGE
myrelease-chart-5755b87555-qjc25 0/1 ContainerCreating 0 1s

NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app=chart,release=myrelease" -o jsonpath=".items[0].metadata.name")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl port-forward $POD_NAME 8080:80

...ng up Helm/example-app [master] ➔ helm upgrade --set image.tag=v1.0.1 myrelease ./chart
Release "myrelease" has been upgraded. Happy Helming!
LAST DEPLOYED: Sun Jun 24 14:59:25 2018
NAMESPACE: default
STATUS: DEPLOYED

RESOURCES:
==> v1/Service
NAME TYPE CLUSTER-IP EXTERNAL-IP PORT(S) AGE
myrelease-chart ClusterIP 10.102.11.62 <none> 80/TCP 1m

==> v1beta1/Deployment
NAME DESIRED CURRENT UP-TO-DATE AVAILABLE AGE
myrelease-chart 1 1 1 0 1m

==> v1/Pod(related)
NAME READY STATUS RESTARTS AGE
myrelease-chart-54f6975cbd-p8bhf 0/1 ContainerCreating 0 1s
myrelease-chart-5755b87555-qjc25 0/1 Terminating 0 1m

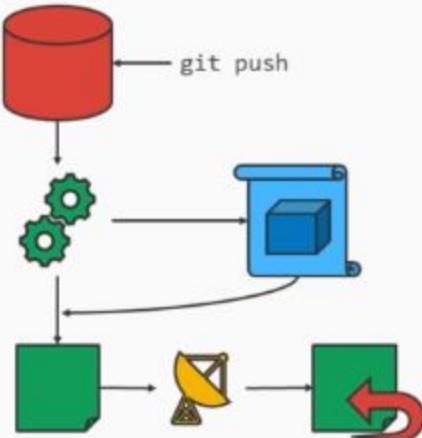
NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app=chart,release=myrelease" -o jsonpath=".items[0].metadata.name")
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl port-forward $POD_NAME 8080:80

...ng up Helm/example-app [master] ➔ helm upgrade --install myrelease ./chart

```

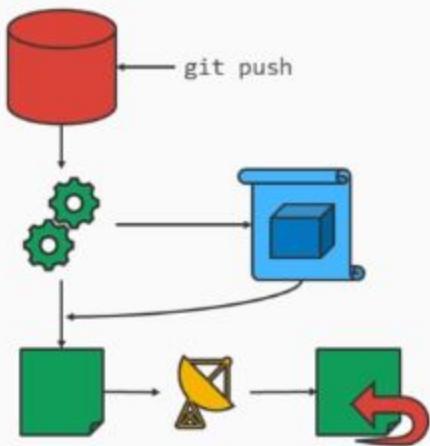
Using Helm in CI

Previously

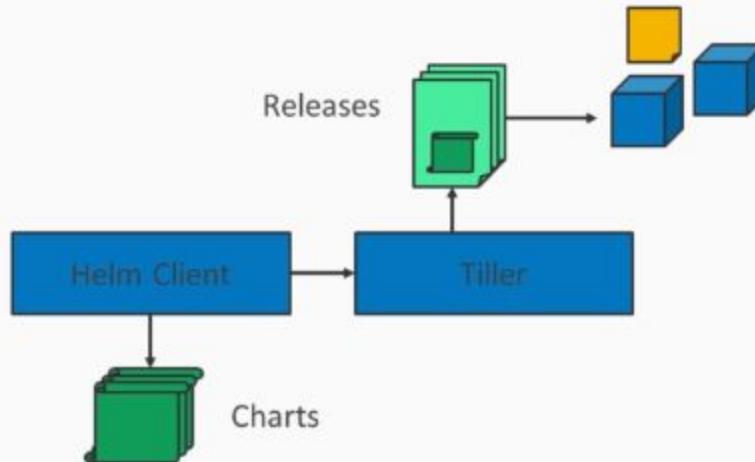


Continuous Delivery
with GitLab

Previously



Continuous Delivery
with GitLab



Kubernetes Package Management
with Helm

Using Helm in GitLab CI

```
deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - >
      kubectl set image deployment/your-app
      app=app:${CI_COMMIT_TAG}
    - kubectl get pods
  only:
    - tags
```

Using Helm in GitLab CI

```
deploy:
  stage: deploy
  image: quay.io/martinhelmich/k8s-deployer:v1.0.0
  environment: development
  script:
    - - >
    -   kubectl set image deployment/your-app
    -     app=app:${CI_COMMIT_TAG}
    + - >
    +   helm upgrade --install
    +     --set image.tag=${CI_COMMIT_TAG}
    +       ${CI_PROJECT_NAME} ./chart
    - kubectl get pods
  only:
    - tags
```

```
...Helm in CI/example-app ➜ helm ci ➤ ls -al
total 11968
drwxr-xr-x  9 mhelmich  staff   288 24 Jun 17:27 .
drwxr-xr-x  3 mhelmich  staff    96 23 Jun 13:21 ..
drwxr-xr-x 14 mhelmich  staff   448 24 Jun 17:35 .git
-rw-r--r--  1 mhelmich  staff   834 24 Jun 17:27 .gitlab-ci.yml
-rw-r--r--  1 mhelmich  staff    78 23 Jun 11:25 Dockerfile
drwxr-xr-x  7 mhelmich  staff   224 24 Jun 17:22 chart
-rw-r--r--  1 mhelmich  staff   298 24 Jun 17:20 deployment.yaml
-rwxr-xr-x  1 mhelmich  staff  6110461 23 Jun 10:19 example-app
-rw-r--r--  1 mhelmich  staff   198 23 Jun 10:18 main.go
```

```
... Helm in CI/example-app ➜ helm ci ➤ ls -al
total 11968
drwxr-xr-x  9 mhelmich  staff      288 24 Jun 17:27 .
drwxr-xr-x  3 mhelmich  staff      96 23 Jun 13:21 ..
drwxr-xr-x 14 mhelmich  staff     448 24 Jun 17:35 .git
-rw-r--r--  1 mhelmich  staff    834 24 Jun 17:27 .gitlab-ci.yml
-rw-r--r--  1 mhelmich  staff     78 23 Jun 11:25 Dockerfile
drwxr-xr-x  7 mhelmich  staff    224 24 Jun 17:22 chart
-rw-r--r--  1 mhelmich  staff   298 24 Jun 17:28 deployment.yaml
-rwxr-xr-x  1 mhelmich  staff  6110461 23 Jun 10:19 example-app
-rw-r--r--  1 mhelmich  staff    198 23 Jun 10:18 main.go
... Helm in CI/example-app ➜ helm ci ➤ ls -al ./chart
total 24
drwxr-xr-x  7 mhelmich  staff    224 24 Jun 17:22 .
drwxr-xr-x  9 mhelmich  staff    288 24 Jun 17:27 ..
-rw-r--r--  1 mhelmich  staff    333 24 Jun 17:22 .helmignore
-rw-r--r--  1 mhelmich  staff     83 24 Jun 17:22 Chart.yaml
drwxr-xr-x  2 mhelmich  staff     64 23 Jun 13:21 charts
drwxr-xr-x  6 mhelmich  staff   192 24 Jun 17:22 templates
-rw-r--r--  1 mhelmich  staff  1689 24 Jun 17:22 values.yaml
... Helm in CI/example-app ➜ helm ci ➤
```

GitLab

Projects Groups Activity Milestones Snippets

This project Search

test

Martin Helmich > test > Pipelines

All 9 Pending 0 Running 1 Finished 8 Branches Tags

Run Pipeline Clear Runner Caches CI Lint

Status	Pipeline	Commit	Stages
running	#9 by  latest	v2.0.0 - 5775740f  Helm CI	  

X

Project Repository Issues Merge Requests

CI / CD

Pipelines Jobs Schedules Charts

Operations Wiki Snippets

GitLab Projects Groups Activity Milestones Snippets ↗

This project Search 🔍 ⌂ ⌂ ⌂ ⌂

test

Martin Helmich > test > Pipelines > #9

passed Pipeline #9 triggered 19 seconds ago by Martin Helmich

Helm CI

⌚ 3 jobs from v2.0.0 (queued for 2 seconds)

→ 5775740f ... 📁

Pipelines

Pipeline Jobs 3

Build Package Deploy

compile dockerbuild deploy

Jobs

Schedules

Charts

Operations

Wiki

Snippets

GitLab Projects Groups Activity Milestones Snippets 🔍 This project Search 🌐

T test

Project Repository Issues Merge Requests 0 CI / CD Pipelines Jobs Schedules Charts Operations Wiki Snippets

```
Starting v2.0.0-debian setup
Downloading artifacts for compile (17)...
Downloading artifacts from coordinator... ok      id=17 responseStatus=200 OK token=V3zm5Un3
$ helm upgrade --install --set image.get=${CI_COMMIT_TAG} ${CI_PROJECT_NAME} ./chart
Release "test" has been upgraded. Happy Helm-ing!
LAST DEPLOYED: Sun Jun 24 15:39:57 2018
NAMESPACE: default
STATUS: DEPLOYED

RESOURCES:
==> v1/Service
NAME        TYPE      CLUSTER-IP     EXTERNAL-IP   PORT(S)    AGE
test-chart  ClusterIP  10.96.212.94  <none>       80/TCP    4m

==> v1beta1/Deployment
NAME        DESIRED  CURRENT  UP-TO-DATE  AVAILABLE  AGE
test-chart  1         1         1           1          4m

==> v1/Pod(related)
NAME                           READY  STATUS    RESTARTS  AGE
test-chart-78c5f58bfb-5v9xj  1/1    Running   0          4m

NOTES:
1. Get the application URL by running these commands:
  export POD_NAME=$(kubectl get pods --namespace default -l "app=chart,release=test" -o
  jsonpath='{.items[0].metadata.name}')
  echo "Visit http://127.0.0.1:8080 to use your application"
  kubectl port-forward $POD_NAME 8080:80

Job succeeded
```

deploy

Duration: 7 seconds
Timeout: 1h (from project)
Runner: gitlab-runner-96bd68669-xrpzf (#1)

Commit 5775740f

Helm CI

Pipeline #9 from v2.0.0

deploy

→ deploy

The End